

分类号: _____


单位代码: 10300

密 级: _____

学 号: 202452200021

南京信息工程大学

学 术 学 位 硕 士 学 位 论 文



论文题目: 基于有界损失函数的噪声鲁棒正-无标记学
习: 理论、算法与深度扩展

申请人姓名: Awasthi Lalit

指导教师: 唐龙 教授

学科名称: 人工智能

研究方向: 机器学习; 弱监督学习; 正-无标记学习; 鲁棒
优化; 有界损失函数设计

培养学院: 人工智能学院

提交时间: 2026 年 05 月 10 日

2026 年 05 月 10 日

Classification Number: _____

University Code: 10300

Secret Level: _____

Student Number: 202452200021

南京信息工程大学

Master's Degree Thesis



Thesis Title: **Noise-Robust Positive-Unlabeled Learning
Through Bounded Loss Functions: Theory,
Algorithms and Deep Extensions**

Applicant: Awasthi Lalit

Supervisor: Professor Tang Long

Major: Artificial Intelligence

Machine Learning; Weak Supervision; Positive-

Research Direction: Unlabeled Learning; Robust Optimization;

Bounded Loss Design

School Affiliated: School of Artificial Intelligence

Submission Time: May, 2026

2026 YY 05 MM 10 DD

原创性声明

本人声明所呈交的论文是我个人在导师指导下进行的研究工作及取得的研究成果。本论文尽我所知，除了文中特别加以标注和致谢的内容外，不包含其他人或其他机构已经发表或撰写过的研究成果，也不包含为获得南京信息工程大学或其他教育机构的学位或证书而使用过的材料。与我一同工作的同志对本研究所做的贡献均已在论文中作了声明并表示谢意。

本人学位论文及涉及相关资料若有不实，愿意承担一切相关的法律责任。

学位论文作者签名： Awasthi Lalit

签字日期： 2026 年 05 月 10 日

关于论文使用授权说明

本人授权南京信息工程大学可以保留并向国家有关部门或机构送交论文的复印件和电子文档；允许论文被查阅和借阅；可以将学位论文的全部或部分内容编入有关数据库进行检索；可以采用影印、缩印或扫描等复制手段保存、汇编本学位论文。本文电子文档的内容和纸质论文的内容相一致。论文的公布（包括刊登）授权南京信息工程大学研究生院办理。

非国家秘密类涉密学位论文在解密后适用本授权书。

☐公开 ☐保密（ 年 月）（保密的学位论文在解密后应遵守此协议）

论文作者签名： Awasthi Lalit

签字日期： 2026 年 05 月 10 日

指导教师签名： Long Long

签字日期： 2026 年 05 月 10 日

Originality Statement

I declare that the paper submitted is my personal research work and research results obtained under the guidance of my supervisor. To the best of my knowledge, this paper does not contain research results that have been published or written by other people or other institutions, except for the content specifically annotated and acknowledged in the article, nor does it contain the results of obtaining a degree or a certificate from Nanjing University of Information Science and Technology or other educational institutions. I have stated and expressed my gratitude in the thesis for the contributions to this research made by the people who worked with me.

If my thesis and related materials are false, I am willing to bear all related legal responsibilities.

Applicant's Signature: Awasthi Lalit

Date: 2026-05-10

Agreement on Authorized Use of Thesis

I authorize Nanjing University of Information Science and Technology to retain and send copies and electronic documents of the thesis to relevant national departments or institutions; allow the thesis to be consulted and borrowed; can compile all or part of the content of the thesis into the relevant database for retrieval; it can be preserved and compiled by photocopy, compact reprint or scanning, etc.. The content of the electronic thesis is consistent with the content of the hardcopy. The publication of the paper is authorized to be handled by the Graduate School of Nanjing University of Information Science and Technology.

This power of attorney shall be applied to non-state secret related thesis after decryption.

☐ Public ☐ Confidential (____ Year ____ Month)

(Confidential thesis should obey this agreement when confidential period is ended)

Applicant's Signature: Awasthi Lalit

Date: 2026-05-10

Supervisor's Signature: Long Long

Date: 2026-05-10

Abstract

Positive-Unlabeled (PU) learning is a critical weakly supervised paradigm for scenarios where only a subset of positive examples is labeled, with applications spanning medical diagnosis, fraud detection, and text classification. However, its practical deployment is severely undermined by label noise, the common occurrence of mislabeled instances within the trusted positive set. Traditional PU methods, including unbiased PU (uPU) and non-negative PU (nnPU), rely on unbounded loss functions that are highly sensitive to such noise, leading to biased risk estimates, unstable training, and degraded generalization in real-world settings.

This thesis addresses this fundamental limitation by proposing a novel robust PU learning framework that integrates bounded loss functions into the PU risk estimation paradigm. We introduce and formalize the use of ramp loss and truncated logistic loss within the nnPU framework, creating a learning objective that inherently limits the influence of any single corrupted label. Our approach is grounded in a comprehensive theoretical analysis, where we derive formal robustness guarantees, proving that the impact of label noise is bounded linearly by the noise rate and the loss function's upper bound.

Empirical validation across synthetic datasets, the UCI Heart Disease dataset, and a PU variant of MNIST demonstrates that our method consistently and significantly outperforms established baselines (uPU and nnPU) under varying noise levels (0–60%). Key metrics including accuracy, ROC-AUC, and PR-AUC show marked improvements, with the ramp loss variant exhibiting exceptional resilience. Furthermore, our framework yields superior training stability, characterized by smoother convergence, lower gradient variance, and reduced sensitivity to hyperparameters.

The primary contributions of this work are threefold: (1) a theoretically robust PU learning framework via bounded loss design, (2) a practical and efficient algorithm suitable for real-world deployment, and (3) extensive empirical evidence of enhanced performance and stability in noisy environments. By bridging the gap between bounded loss theory and PU learning practice, this research advances the state of the art in robust weakly supervised learning, offering a principled and effective solution for building reliable models from imperfect, real-world data.

Keywords: Positive-unlabeled learning; label noise robustness; bounded loss functions; weak supervision; risk estimation

摘要

正例-无标签学习 (Positive-Unlabeled Learning) 是一种重要的弱监督学习范式, 适用于仅有一部分正例被标记的场景, 广泛应用于疾病诊断、欺诈检测和文本分类等领域。然而, 其实际部署常受到标签噪声的严重影响——即被信任的正例集合中常存在错误标记的样本。传统的 PU 学习方法, 包括无偏 PU (uPU) 和非负 PU (nnPU), 依赖于无界损失函数, 对标签噪声极为敏感, 导致风险估计偏差、训练不稳定以及在真实场景下泛化性能下降。

本论文针对这一根本性局限, 提出了一种新颖的鲁棒 PU 学习框架, 将有界损失函数集成到 PU 风险估计范式中。我们引入并在 nnPU 框架中形式化使用了斜坡损失和截断逻辑损失, 构建了一个能够从根本上限制单个错误标签影响的学习目标。我们的方法建立在全面的理论分析基础上, 推导出形式化的鲁棒性保证, 证明了标签噪声的影响受到噪声率与损失函数上界的线性约束。

在合成数据集、UCI 心脏病数据集以及 MNIST 的 PU 变体上的实验验证表明, 在不同噪声水平 (0%–60%) 下, 我们的方法始终显著优于传统基线方法 (uPU 和 nnPU)。关键指标——包括准确率、ROC-AUC 和 PR-AUC——均显示出显著提升, 其中斜坡损失变体表现出卓越的抗噪性。此外, 我们的框架实现了更优的训练稳定性, 表现为更平滑的收敛、更低的梯度方差以及对超参数更低的敏感性。

本研究的主要贡献有三方面: (1) 提出了一种基于有界损失设计的理论鲁棒的 PU 学习框架; (2) 提出了一种适用于实际部署的实用高效算法; (3) 提供了在噪声环境下性能与稳定性显著提升的广泛实证证据。通过弥合有界损失理论与 PU 学习实践之间的差距, 本研究推动了鲁棒弱监督学习的前沿发展, 为从真实世界不完美数据中构建可靠模型提供了一个原则性强且有效的解决方案。

关键词: 正例-无标签学习; 标签噪声鲁棒性; 有界损失函数; 弱监督; 风险估计

Table of Contents

原创性声明.....	I
关于论文使用授权说明	I
Originality Statement.....	II
Agreement on Authorized Use of Thesis	II
Abstract.....	III
摘要	IV
Abbreviations.....	VIII
Chapter I INTRODUCTION	1
1.1 Introduction	1
1.2 Background of Positive-Unlabeled Learning	3
1.3 Problem Statement: Label Noise in PU Learning	5
1.4 Research Objectives	7
1.5 Scope and Limitations	7
1.6 Significance and Contributions of the Study.....	9
1.7 Thesis Structure	11
Chapter II LITERATURE REVIEW	13
2.1 Overview of Positive-Unlabeled Learning.....	13
2.1.1 Foundational PU Frameworks (uPU, nnPU).....	14
2.1.2 Challenges in Real-World PU Applications	15
2.2 Label Noise in Machine Learning	18
2.2.1 Types and Impact of Label Noise.....	19
2.2.2 Noise-Robust Learning Methods	22
2.3 Bounded Loss Functions in Robust Learning	25
2.3.1 Ramp Loss and Truncated Logistic Loss	26
2.3.2 Applications in Supervised and Weakly-Supervised Settings.....	29
2.4 Research Gaps in Noise-Robust PU Learning.....	32
2.4.1 Theoretical Gap: Absence of Formal Robustness Guarantees for PU Learning Under Label Noise	33
2.4.2 Methodological Gap.....	33
2.4.3 Empirical Gap: Limited Validation Under Realistic Noise Conditions	34

2.4.4 Integrative Gap: Disconnected Research Trajectories	35
2.4.5 Bridging the Gaps: The Contribution of This Research.....	36
2.5 Chapter Summary	37
Chapter III METHODOLOGY	39
3.1 Introduction	39
3.2 Preliminaries and Problem Formulation.....	40
3.2.1 PU Learning Setup and Risk Estimation.....	40
3.2.2 Label Noise Model	41
3.3 Proposed Framework: Bounded Loss PU Learning	43
3.3.1 Motivation for Bounded Loss Functions.....	43
3.3.2 Ramp Loss and Truncated Logistic Loss Formulations	45
3.3.3 Robust PU Risk Estimator	47
3.4 Theoretical Analysis	48
3.4.1 Robustness Guarantees Under Label Noise	48
3.4.2 Bounded Loss Properties and Convergence.....	50
3.5 Algorithm Description	51
3.5.1 Training Procedure (Algorithm 1).....	51
3.5.2 Implementation Details	53
3.6 Evaluation Metrics.....	53
3.6.1 Accuracy, Precision, Recall, F1-Score	53
3.6.2 ROC-AUC and PR-AUC	54
3.6.3 Training Stability Metrics	54
3.7 Chapter Summary	55
Chapter IV EXPERIMENT RESULTS AND DISCUSSIONS	57
4.1 Experimental Setup	57
4.1.1 Dataset Description (Synthetic, UCI Heart, MNIST-PU)	57
4.1.2 Hardware and Software Configurations	58
4.1.3 Noise Injection Protocol.....	59
4.2 Performance Under Varying Noise Levels	60
4.2.1 Quantitative Results on Synthetic Data.....	60
4.2.2 Results on Real-World Datasets (UCI, MNIST).....	61
4.3 Comparison with Baseline Models.....	63
4.3.1 uPU vs. nnPU vs. Proposed Method	63
4.3.2 Performance Trends with Increasing Noise	64

4.4 Training Behavior and Convergence Analysis	65
4.4.1 Loss Curves and Optimization Stability.....	65
4.4.2 Sensitivity to Hyperparameters	66
4.5 Ablation Studies.....	68
4.5.1 Effect of Bounded Loss vs. Rectification.....	68
4.5.2 Component-Wise Contribution Analysis.....	69
4.6 Discussion on Robustness and Generalization	70
4.6.1 Interpretation of Theoretical Bounds	70
4.6.2 Practical Implications for Noisy PU Scenarios	71
4.7 Chapter Summary	72
Chapter V CONCLUSIONS	73
References.....	76
Research Output Achieved During Master’s Degree Studies	82
Acknowledgement	83

Abbreviations

Abbreviation	-	Full Form
AUC	-	Area Under the Curve
CPU	-	Central Processing Unit
FN	-	False Negatives
FP	-	False Positives
FPR	-	False Positive Rate
GPU	-	Graphics Processing Unit
i.i.d.	-	Independent and Identically Distributed
MAE	-	Mean Absolute Error
MIL	-	Multiple-Instance Learning
MNIST	-	Modified National Institute of Standards and Technology (dataset)
nnPU	-	Non-negative Positive-Unlabeled
PR AUC	-	Precision-Recall Area Under the Curve
PU	-	Positive-Unlabeled
RAM	-	Random Access Memory
ReLU	-	Rectified Linear Unit
ROC AUC	-	Receiver Operating Characteristic Area Under the Curve
SCAR	-	Selected Completely at Random
SGD	-	Stochastic Gradient Descent
SNR	-	Signal-to-Noise Ratio
SSD	-	Solid State Drive
TN	-	True Negatives
TP	-	True Positives
TPR	-	True Positive Rate
UCI	-	University of California, Irvine (dataset source)
uPU	-	Unbiased Positive-Unlabeled

Chapter I INTRODUCTION

1.1 Introduction

The practical deployment of artificial intelligence systems often confronts the reality of imperfect data, making the development of learning algorithms resilient to such deficiencies a critical research frontier [1]. Among the various paradigms addressing this challenge, Positive-Unlabeled (PU) learning has emerged as a critical framework for situations where only a subset of positive examples is identified, while the remainder of the data lacks explicit labeling [2,3]. This setting mirrors numerous real-world scenarios from disease screening and fraud detection to information retrieval and rare event prediction where annotating negative instances is economically prohibitive, logistically complex, or fundamentally ambiguous. PU learning provides a principled methodology to harness both labeled positives and unlabeled data, offering a pathway to develop classifiers when fully supervised training is infeasible [4].

However, the practical deployment of PU learning is frequently undermined by a pervasive yet often overlooked problem: label noise. In many applications, the assumption that the labeled positive set is clean and reliable is untenable. Data collection processes are prone to errors; subjective human annotation introduces inconsistencies; and automated labeling systems can produce inaccuracies [5]. Consequently, the so-called “positive” set may be contaminated with mislabeled examples, instances that actually belong to the negative class [6]. This corruption introduces a distributional shift between the empirical data used for training and the true underlying distribution, leading to biased risk estimation, degraded model performance, and reduced generalization capability. The problem is particularly acute in PU learning because the entire learning signal heavily relies on a small, trusted set of positives [7]. When this foundation is compromised, the resulting classifier can deviate significantly from the optimal decision boundary.

Traditional PU learning methods, such as the unbiased PU (uPU) and non-negative PU (nnPU) frameworks, have made significant strides in risk estimation from positive and unlabeled data [8]. Yet, these approaches predominantly employ standard surrogate loss functions like logistic loss or hinge loss which are inherently unbounded [9,10]. This unbounded nature means that a single mislabeled or outlier sample can exert disproportionately large influence on the gradient during training, destabilizing the optimization process and steering the model toward suboptimal solutions.

This thesis confronts this critical limitation by proposing a robust PU learning framework that integrates bounded loss functions, specifically, ramp loss and truncated logistic loss into the PU risk estimation paradigm. The core innovation lies in capping the maximum contribution any single example can make to the overall loss, thereby inherently dampening the influence of mislabeled positives. This approach provides a natural and principled mechanism to enhance noise robustness without requiring additional clean validation data, complex noise modeling, or architectural overhauls.

The proposed framework is grounded in a rigorous theoretical analysis, where we reformulate the standard PU risk estimator by integrating bounded surrogate loss functions and subsequently derive formal robustness guarantees under conditions of label contamination. We analytically demonstrate that the degradation in the estimator’s performance due to corrupted positive labels is linearly bounded by the product of the noise rate and the maximum attainable loss value, thereby providing provable assurances of stability and predictable performance decay. This theoretical foundation establishes boundedness not as an incidental trait but as a fundamental, quantifiable property that directly mitigates the error propagation inherent in the PU risk decomposition.

Algorithmically, we adapt the established nnPU (non-negative PU) training procedure to seamlessly incorporate bounded loss functions, ensuring that risk estimates remain non-negative while preserving optimization efficiency and convergence properties. A comprehensive empirical evaluation across controlled synthetic data and challenging real-world datasets including the UCI Heart Disease repository and a constructed PU variant of MNIST systematically validates the framework’s practical superiority. Under varying and often high levels of simulated label noise, our proposed method consistently and significantly outperforms established baselines across all key performance metrics: classification accuracy, ROC-AUC, and precision-recall AUC. Notably, the ramp loss variant demonstrates exceptional resilience, maintaining robust performance where traditional unbounded losses exhibit severe degradation, thereby confirming the critical role of boundedness in achieving reliable learning from noisy positive labels.

This research introduces a robust and theoretically grounded framework that integrates bounded loss functions into positive-unlabeled learning. We demonstrate that boundedness is a critical property for mitigating label noise, offering a principled solution to a core challenge in weakly supervised learning. The findings advance the potential for reliable AI in domains such as medical diagnostics and cybersecurity, where data imperfections are common but system reliability is essential.

1.2 Background of Positive-Unlabeled Learning

Positive-Unlabeled learning represents a significant departure from traditional supervised learning paradigms by addressing scenarios where explicit negative labels are unavailable or impractical to obtain [11]. This learning framework operates under the realistic constraint that only a subset of positive examples is identified, while the remaining data potentially containing both positive and negative instances remains unlabeled. The origins of PU learning can be traced to early work in text classification and biomedical informatics [12], where researchers confronted the challenge of building classifiers from partially labeled corpora or biological datasets with confirmed positives but uncertain negatives. Over time, the paradigm has evolved from ad hoc solutions to a mature subfield with formal theoretical foundations, driven by the recognition that many real-world problems such as disease screening, rare event detection, and information retrieval, inherently possess this asymmetric label structure. Consequently, PU learning has established itself not merely as a niche technique, but as an essential methodological toolkit for weakly supervised learning [13,14], bridging the gap between fully supervised classification and completely unsupervised discovery.

The foundational assumption underlying PU learning is that the labeled positive examples are randomly sampled from the true positive distribution, a premise known as the Selected Completely at Random (SCAR) assumption [15]. This assumption allows for the derivation of unbiased risk estimators that compensate for the missing negative labels by leveraging the statistical properties of the unlabeled data. Formally, PU learning aims to train a binary classifier $f: \mathcal{X} \rightarrow \mathbb{R}$ by minimizing an expected risk expressed in terms of the observable distributions: the labeled positive distribution $p(x | Y = +1)$ and the marginal distribution $p(x)$ of the unlabeled data [16].

Within this framework, the unlabeled data is understood as a mixture of hidden positive and negative examples, with the class prior $\pi = p(Y = +1)$ playing a central role in reweighting the contributions of each component. The resulting risk estimators such as the unbiased PU (uPU) and non-negative PU (nnPU) risks enable model training without access to explicitly labeled negatives, thereby extending supervised classification to realistic settings where negative labels are unavailable or prohibitively costly to obtain [17].

Over the past decade, PU learning has evolved along two primary methodological branches:

1) Risk Estimation Approaches: These methods reformulate the supervised classification risk using only positive and unlabeled data. The unbiased PU (uPU) estimator, introduced by

du Plessis et al. [18], provides an unbiased risk estimate but suffers from high variance, especially with flexible models. To address this, the non-negative PU (nnPU) estimator applies a rectification mechanism to ensure non-negative empirical risks, yielding more stable training.

2) Two-Step Strategies: These involve first identifying reliable negative examples from the unlabeled set (e.g., via spy techniques, clustering, or iterative refinement), then training a standard classifier on the augmented labeled set [7]. While intuitive, these methods are sensitive to the initial identification step and may propagate errors.

The diagram below (Figure 1-1) illustrates the conceptual framework and methodological evolution of PU learning:




Figure 1-1 Conceptual overview and methodological evolution of positive-unlabeled learning (Source: Created by author)

PU learning has established itself as a versatile and successful paradigm with compelling applications across diverse, high-impact domains. In medical diagnosis and screening, it enables predictive models to be trained using only confirmed disease cases such as biopsy-verified tumors or PCR-positive infections as positives, combined with a general patient population where disease status is unknown or unverified, thereby circumventing the often impractical requirement for exhaustive negative confirmations [4]. In financial surveillance

and fraud detection, the framework allows security systems to learn from a limited set of known fraudulent transactions while leveraging the vast volume of unlabeled daily transactions to identify novel and evolving fraudulent patterns without costly manual review of every legitimate transaction [19,20].

Similarly, in text mining and information retrieval, PU learning effectively trains classifiers from a curated set of relevant documents, articles on a specific topic or reviews expressing a particular sentiment and a massive, unlabeled general corpus, significantly reducing annotation overhead [12]. While these real-world applications powerfully demonstrate the practical utility and scalability of the PU framework, they also starkly expose its intrinsic vulnerability to label noise. This challenge becomes particularly acute and consequential when the presumed-trusted positive set itself contains mislabeled examples such as a misdiagnosed patient, a mischaracterized fraudulent transaction, or an incorrectly categorized document as even a small proportion of such errors can corrupt the risk estimation process and severely degrade model reliability in deployment.

The growing recognition of label noise as a critical bottleneck in PU learning has spurred research into robust alternatives [21]. However, most existing solutions either assume access to clean validation data or focus on modeling biases in the unlabeled set, leaving the problem of noisy positive labels inadequately addressed [22]. This gap motivates the development of a robust PU learning framework that systematically mitigates the impact of label corruption through principled loss function design, an approach that forms the core contribution of this thesis.

1.3 Problem Statement: Label Noise in PU Learning

The practical application of Positive-Unlabeled learning is often compromised by label noise within the labeled positive set. In practice, labeled positives are frequently impure, containing mislabeled instances due to annotation errors, automated labeling flaws, or data integration issues. This contamination breaks the core SCAR assumption, creates a mismatch between the observed and true positive distributions, and ultimately results in biased risk estimators and poorer model performance.

The problem is especially pronounced in PU learning for two interconnected reasons:

1) Heavy Reliance on a Small Positive Set: PU learning algorithms anchor their risk estimation and calibration on the assumption that the labeled positives are trustworthy. When this set is corrupted, the entire learning signal becomes distorted, causing the classifier to learn incorrect decision boundaries.

2) Amplification by Unbounded Loss Functions: Traditional PU methods rely on unbounded surrogate loss functions (e.g., logistic, exponential). These functions assign disproportionately large penalties to mislabeled examples, especially those far from the decision boundary causing unstable gradients, overfitting to noise, and poor generalization.

Table 1-1 summarizes the core challenges posed by label noise in PU learning and their implications for model performance and deployment:

Table 1-1 Key challenges and implications of label noise in PU learning

Challenge	Description	Impact on PU Learning	Practical Consequence
Violation of SCAR Assumption	Labeled positives are not a random sample from the true positive class due to noise.	Biased risk estimators; model converges to suboptimal parameters.	Reduced classification accuracy and reliability.
Unbounded Loss Sensitivity	Standard loss functions (logistic, hinge) penalize outliers and mislabeled examples excessively.	Training instability; gradient explosions; overfitting to noisy samples.	Poor generalization; high variance in performance.
Noise Amplification in Risk Estimation	Noise in positives propagates through risk terms, especially in the negative component estimation.	Over- or under-estimation of the negative class risk.	Increased false positives or false negatives.
Lack of Explicit Noise Modeling	Most PU methods ignore label noise or assume clean positives.	No inherent mechanism to detect or correct mislabeled positives.	Models fail in real-world noisy settings.
Class Prior Distortion	Estimated class prior π becomes inaccurate when positives are contaminated.	Incorrect weighting of positive and negative risks.	Degraded calibration and decision thresholds.

(Source: Created by author)

As illustrated in Table 1-1, the absence of explicit noise-handling mechanisms in conventional PU learning renders existing methods including uPU and nnPU, particularly fragile in real-world applications where data imperfection is the norm rather than the exception. While some recent studies have begun to address label noise in fully supervised or semi-supervised settings, the unique structure of PU learning with its complete absence of verified negatives requires a specialized approach that does not rely on clean validation data or complex noise correction pipelines.

Therefore, the central problem addressed in this thesis is: How can we design a robust PU learning framework that maintains stable performance and reliable risk estimation when the labeled positive set contains mislabeled or outlier instances? We propose that the answer lies not in complex noise modeling or auxiliary data, but in a fundamental redesign of the loss function itself.

1.4 Research Objectives

The primary aim of this research is to develop a robust and theoretically grounded framework for Positive-Unlabeled learning that remains effective in the presence of label noise. Specifically, this thesis pursues the following interconnected objectives:

1) To formulate a robust PU learning framework that systematically integrates bounded loss functions (ramp loss and truncated logistic loss) into the non-negative PU (nnPU) risk estimation paradigm, thereby limiting the undue influence of mislabeled positive examples during training.

2) To provide a theoretical foundation for the proposed framework by deriving formal robustness guarantees under label noise. This includes analyzing risk consistency, bounding the impact of corrupted labels, and establishing convergence properties under bounded loss functions.

3) To develop an efficient algorithmic implementation of the proposed bounded-loss PU learning method, complete with optimization details, hyperparameter guidelines, and a reproducible training procedure suitable for both synthetic and real-world datasets.

4) To conduct comprehensive empirical validation of the proposed framework across multiple datasets and noise regimes, benchmarking its performance against established PU baselines (uPU and nnPU) in terms of classification accuracy, ROC-AUC, PR-AUC, and training stability.

5) To explore the practical implications and extensibility of the framework by discussing its integration into real-world applications, identifying limitations, and proposing viable directions for future research including extensions to deep learning, non-uniform noise models, and semi-supervised settings.

1.5 Scope and Limitations

This research focuses specifically on enhancing the robustness of binary Positive-Unlabeled learning to label noise within the positive set. The scope is deliberately bounded to ensure depth and rigor in addressing this well-defined challenge. The work operates under several key assumptions and acknowledges inherent limitations, which are summarized in Table 1-2 and further illustrated in Figure 1-2. This study aims to provide a thorough and replicable investigation into the core problem, establishing a solid foundation upon which future work can expand to address the noted limitations.

Table 1-2 Scope and key limitations of the research

Aspect	Scope	Limitations
Learning Paradigm	Binary Positive-Unlabeled (PU) classification.	Does not extend to multi-class or multi-label PU settings.
Noise Type	Uniform random label noise in the positive set.	Does not address instance-dependent, structured, or adversarial noise patterns.
Data Requirements	Requires labeled positives and unlabeled mixed samples; class prior (π) known or estimable.	Assumes the SCAR condition holds for clean positives; does not leverage additional clean validation data.
Loss Functions	Focus on bounded surrogates: ramp loss and truncated logistic loss.	Other bounded losses (e.g., Savage, 0-1) are not exhaustively explored.
Theoretical Analysis	Provides robustness bounds under uniform noise and Lipschitz continuity.	Does not characterize exact error rates under complex noise distributions.
Empirical Validation	Evaluated on synthetic, medical (UCI Heart), and image (MNIST-PU) datasets.	Not tested on extremely high-dimensional or sequential data domains.
Computational Focus	Algorithmic efficiency and training stability are considered.	Real-time deployment on edge devices or large-scale distributed systems is not evaluated.

(Source: Created by author)

These delineated boundaries serve to maintain analytical tractability while ensuring that the proposed methodologies are rigorously evaluated within their intended context. By explicitly acknowledging these constraints, the research maintains transparency regarding the generalizability of its findings. The controlled scope enables a deeper investigation into the fundamental mechanisms of noise-robust PU learning without the confounding effects of additional complexities such as instance-dependent noise or multi-class scenarios. Consequently, the conclusions drawn from this work should be interpreted within these specified boundaries, and practitioners are advised to consider these limitations when applying the proposed methods to domains that fall outside the stated scope.

While this research makes significant strides in robustifying PU learning against label noise, it is important to acknowledge its defined boundaries and the specific context of its contributions. The proposed framework analytically assumes a uniform random noise model, which, while providing essential tractability for deriving formal guarantees, may not fully capture the more complex structured, class-conditional, or instance-dependent noise patterns often encountered in real-world annotation processes. Furthermore, the methodology presumes that the class prior π is either known or can be accurately estimated from the data, an assumption that may become tenuous in severely imbalanced scenarios or under conditions of extreme label contamination, potentially impacting the calibrated weighting of risk components.

Consequently, the empirical evaluation, though comprehensive and systematic across the chosen synthetic, medical, and image benchmark datasets, does not cover the full spectrum of potential application domains. Performance in rich, high-dimensional modalities such as text, audio, or sequential video data remains an open question for future investigation. Future research directions may include relaxing the SCAR assumption, extending to multi-class problems, and investigating more complex noise structures that better reflect real-world data conditions.

These limitations, however, are not endpoints but rather foundations for future work. They highlight clear pathways for extension, such as adaptive noise modeling, integration with deep representation learning, and application to more complex data modalities. By transparently delineating its scope and constraints, this thesis provides a solid and replicable basis for subsequent research in noise-robust weakly supervised learning.




Figure 1-2 Visual scope map of the research (Source: Created by author)

1.6 Significance and Contributions of the Study

This research makes several significant and interrelated contributions to the field of weakly supervised machine learning, with particular emphasis on enhancing the robustness and practical applicability of Positive-Unlabeled learning systems in real-world, noisy environments. By tackling the underexplored yet pervasive problem of label noise within the

positive set, this study provides a crucial bridge between theoretical soundness and practical deployability. The study addresses a critical gap in existing literature by systematically integrating bounded loss functions into the PU learning paradigm, an approach that offers both theoretical elegance and empirical effectiveness while ensuring resilience against corrupted annotations. The key contributions of this thesis are threefold:

1) **Theoretical Advancement:** The research provides a formal theoretical foundation for robust PU learning under label noise. By reformulating the PU risk estimator with bounded loss functions and deriving explicit robustness bounds, this work demonstrates that the impact of corrupted labels is linearly bounded by both the noise rate and the maximum loss value. This contribution extends the theoretical understanding of noise tolerance in weakly supervised settings and establishes bounded loss design as a principled strategy for achieving robustness.

2) **Methodological Innovation:** This thesis introduces a novel, practical framework that enhances the nnPU learning algorithm through the integration of bounded loss functions—specifically, ramp loss and truncated logistic loss. The proposed method is algorithmically simple, computationally efficient, and requires no additional clean validation data or complex noise correction pipelines. It represents a significant step forward in developing "noise-aware" PU learning systems that maintain performance stability without sacrificing optimization efficiency.

3) **Empirical Validation:** Through comprehensive experiments on synthetic and real-world datasets, this research provides robust empirical evidence that bounded loss functions consistently outperform traditional unbounded alternatives under varying noise conditions. The systematic evaluation across multiple metrics (accuracy, ROC-AUC, PR-AUC) and noise levels (0-60%) demonstrates the practical utility of the proposed approach in domains where label noise is unavoidable, such as medical diagnosis and fraud detection.

Beyond these direct contributions, this study has important implications for the broader machine learning community. It highlights the critical role of loss function design in achieving robustness, challenges the prevailing reliance on unbounded surrogates in weakly supervised settings, and provides a replicable framework that can be extended to related problems such as semi-supervised learning, noisy label learning, and domain adaptation. The findings advocate for a paradigm shift toward more deliberate, theoretically-informed loss function selection in practical machine learning systems.

1.7 Thesis Structure

This thesis is organized into five core chapters that logically progress from establishing the research problem to presenting the solution, its validation, and its implications. The structure is designed to provide a coherent and comprehensive narrative of the research journey, ensuring both theoretical depth and practical relevance.

Chapter 1: Introduction

This chapter has established the foundational context for the research. It introduced the critical problem of label noise in Positive-Unlabeled learning, provided a background on the evolution of PU learning frameworks and their vulnerabilities, articulated a clear problem statement highlighting the limitations of existing approaches, and outlined the research objectives and questions guiding this work. It has also defined the scope of the study and summarized its significant contributions to the field of weakly supervised machine learning.

Chapter 2: Literature Review

This chapter will provide a comprehensive and critical survey of the existing body of knowledge relevant to this thesis. The taxonomy of Positive-Unlabeled learning methods will be thoroughly examined, including risk estimation approaches and two-step strategies. The landscape of label noise robustness techniques in machine learning will be methodically reviewed and categorized, with particular attention to bounded loss functions and their applications. A focused analysis of applications in medical diagnosis, fraud detection, and text mining will be presented, culminating in a precise identification of the research gaps that this thesis aims to fill.

Chapter 3: Methodology

This chapter presents the core methodological framework of the research. It details the formulation of robust PU learning through bounded loss functions, including the mathematical derivations of ramp loss and truncated logistic loss within the nnPU risk estimation paradigm. The theoretical analysis section provides robustness guarantees and convergence properties under label noise. The chapter concludes with a structured algorithmic description of the training procedure, implementation details, and evaluation metrics used throughout the experimental validation.

Chapter 4: Experiment Results and Discussions

This chapter presents comprehensive empirical validation of the proposed framework. It details the experimental setup, including datasets, noise injection protocols, and baseline comparisons. Quantitative results across synthetic and real-world datasets are analyzed under

varying noise conditions, with particular attention to performance metrics and training stability. The chapter includes ablation studies, sensitivity analyses, and graphical visualizations that collectively demonstrate the superiority of the bounded loss approach over traditional PU methods.

Chapter 5: Conclusions and Future Work

The final chapter synthesizes the entire research endeavor. It summarizes the key findings, restates the thesis's contributions in the context of the field, and reflects on the broader implications of the work for the future of robust machine learning. The thesis will conclude by candidly addressing the limitations of the current study and proposing concrete, promising directions for future research, including extensions to deep learning architectures, adaptive noise modeling, and integration with semi-supervised learning techniques.

The document will be completed with supporting sections including **Research Output Achieved During Master's Degree Studies**, a comprehensive **References** list, and an appreciation section on **Acknowledgements**,

Chapter II LITERATURE REVIEW

2.1 Overview of Positive-Unlabeled Learning

Positive-Unlabeled (PU) learning has emerged as a vital subfield within weakly supervised machine learning, addressing the pervasive real-world scenario where only a subset of positive examples is explicitly labeled, while the remaining data exists without class identifiers [23]. This paradigm fundamentally challenges the traditional binary classification framework by operating under the constraint that negative labels are either unavailable, too costly to obtain, or inherently ambiguous. The genesis of PU learning can be traced to pioneering work in text classification and information retrieval [24], where researchers confronted the practical challenge of building classifiers from document collections where only relevant documents (positives) could be reliably identified, while the remainder constituted an unlabeled mixture of both relevant and irrelevant material.

Formally, PU learning operates on two distinct data subsets: a labeled positive set $\mathcal{P} = \{x_i^p\}_{i=1}^{n_p}$ drawn independently from the true positive distribution $p(x | Y = +1)$, and an unlabeled set $\mathcal{U} = \{x_i^u\}_{i=1}^{n_u}$ drawn from the marginal distribution $p(x) = \pi p(x | Y = +1) + (1 - \pi)p(x | Y = -1)$, where $\pi = P(Y = +1)$ represents the class prior. The central learning objective remains consistent with standard classification: to train a classifier $f: \mathcal{X} \rightarrow \mathbb{R}$ that minimizes the expected risk $R(f) = \mathbb{E}_{(x,y)}[\ell(y, f(x))]$. However, the absence of explicit negative labels necessitates innovative risk estimation strategies that reconstruct the complete supervised risk using only observable components [25].

The theoretical foundation of modern PU learning rests upon the Selected Completely at Random (SCAR) assumption, which posits that labeled positives constitute a random sample from the true positive distribution [26]. This critical assumption enables the derivation of unbiased risk estimators and ensures the identifiability of the learning problem. Under SCAR, the labeling mechanism is independent of the features, meaning that the probability of a positive example being labeled depends only on the class, not on its specific attributes. While this assumption has been debated and relaxed in subsequent work, it remains a cornerstone of most practical PU learning algorithms. PU learning methodologies have evolved along two primary trajectories: risk estimation approaches that directly reformulate the classification risk using statistical properties of the observed data, and two-step strategies that first identify

reliable negatives from the unlabeled pool before applying standard supervised learning [8,27]. The risk estimation paradigm has gained prominence due to its theoretical elegance and direct optimization characteristics, forming the foundation upon which this research builds.

2.1.1 Foundational PU Frameworks (uPU, nnPU)

The development of formal risk estimation frameworks represents the most significant theoretical advancement in PU learning [28]. These approaches circumvent the need for explicit negative identification by algebraically reconstructing the complete classification risk from observable components.

The Unbiased PU (uPU) learning framework, introduced by du Plessis et al. [18], represents a landmark contribution that established PU learning as a theoretically grounded discipline. The uPU risk estimator is derived by expressing the negative class expectation in terms of the marginal and positive distributions [29]:

$$R_{uPU}(f) = \pi \cdot \mathbb{E}_{x \sim p(x|Y=+1)}[\ell(+1, f(x))] + \mathbb{E}_{x \sim p(x)}[\ell(-1, f(x))] - \pi \cdot \mathbb{E}_{x \sim p(x|Y=+1)}[\ell(-1, f(x))] \quad (2-1)$$

This formulation possesses the crucial property of being unbiased with respect to the true supervised risk, meaning that $\mathbb{E}[R_{uPU}(f)] = R(f)$ under the SCAR assumption. The estimator works by treating the unlabeled data as a weighted combination of positive and negative samples, then subtracting the expected contribution of positives to isolate the negative component. Despite its theoretical appeal, uPU suffers from practical limitations, most notably high variance in empirical settings. When the classifier is flexible or the loss function is unbounded, the subtraction of large opposing terms can lead to unstable estimates and potentially negative empirical risks, which lack meaningful interpretation in the classification context.

To address these limitations, Kiryo et al. introduced the Non-negative PU (nnPU) learning framework [28], which modifies the uPU estimator by imposing a non-negativity constraint:

$$R_{nnPU}(f) = \pi \cdot \mathbb{E}_{x \sim p(x|Y=+1)}[\ell(+1, f(x))] + \max(0, \mathbb{E}_{x \sim p(x)}[\ell(-1, f(x))] - \pi \cdot \mathbb{E}_{x \sim p(x|Y=+1)}[\ell(-1, f(x))]) \quad (2-2)$$

The rectification operator $\max(0, \cdot)$ ensures that the empirical risk remains non-negative throughout optimization, stabilizing training and preventing the pathological behavior observed in uPU [30]. This great modification transforms the unbiased estimator into

a consistent one that converges to the true risk while maintaining practical stability. The nnPU framework has become the de facto standard for modern PU learning applications, demonstrating superior performance across diverse domains while preserving theoretical rigor.

Table 2-1 Comparative analysis of foundational PU learning frameworks

Framework	Key Formulation	Theoretical Property	Practical Advantages	Primary Limitations
Unbiased PU (uPU)	$R_{uPU} = \pi E p[\ell+] + E u[\ell-] - \pi E p[\ell-]$	Unbiased risk estimator	Theoretical elegance; consistent under SCAR	High variance; negative empirical risks; unstable training
Non-negative PU (nnPU)	$R_{nnPU} = \pi E p[\ell+] + \max(0, E u[\ell-] - \pi E p[\ell-])$	Consistent risk estimator	Stable optimization; non-negative risks; robust performance	Potential underestimation bias; sensitivity to class prior estimation
Two-step Methods	1. Identify reliable negatives 2. Train supervised classifier	Heuristic foundation	Intuitive; leverages mature supervised methods	Error propagation; sensitive to initial step; lacks theoretical guarantees

(Source: Created by author)

As summarized in Table 2-1, both uPU and nnPU operate within the risk estimation paradigm but address different aspects of the learning problem. While uPU prioritizes theoretical purity through unbiased estimation, nnPU emphasizes practical utility through stabilized optimization. The evolution from uPU to nnPU reflects a broader trend in machine learning toward methods that balance theoretical soundness with empirical effectiveness, a balance that this research extends by incorporating robustness considerations into the nnPU framework.

2.1.2 Challenges in Real-World PU Applications

Despite significant theoretical advancements, the practical deployment of PU learning systems continues to face substantial challenges that limit their effectiveness in real-world applications. These challenges emerge from the complex interplay between theoretical assumptions and practical constraints, often revealing gaps between idealized formulations and operational realities.

Label Quality and Noise Contamination represent perhaps the most critical challenge in practical PU learning [31–33]. The SCAR assumption presupposes that labeled positives are both accurate and representative of the true positive class [34]. In reality, data collection processes introduce multiple forms of noise: mislabeling errors where negative instances are incorrectly annotated as positive; annotation inconsistencies arising from subjective human

judgment; and systematic biases in data acquisition that violate the random sampling assumption. As documented in medical diagnostics research, label error rates in specialized domains can exceed 20%, severely compromising model performance [35]. This noise contamination is particularly damaging in PU learning because the entire risk estimation framework relies heavily on the integrity of the positive set, a single corrupted component that propagates errors throughout the learning process.

Class Prior Estimation Uncertainty presents another fundamental challenge. Most PU risk estimators, including both uPU and nnPU, require knowledge of the class prior π , the proportion of positives in the entire population [36]. In practice, π is rarely known with certainty and must be estimated from the unlabeled data. While several estimation methods have been proposed, including mixture proportion estimation and envelope approaches, these estimators exhibit sensitivity to distributional assumptions and can introduce substantial errors when assumptions are violated. An inaccurate π estimate systematically biases risk estimation, leading to either overemphasis or underemphasis of the negative component and consequent degradation in classifier performance [37].

Distributional Shift and Covariate Drift further complicate real-world deployment. PU learning algorithms typically assume that both labeled positives and unlabeled data are drawn from stationary distributions that remain consistent between training and deployment [38]. In dynamic environments such as fraud detection systems or evolving medical diagnostics, this assumption frequently fails, new fraud patterns emerge, disease manifestations evolve, and user behaviors change over time. Such distributional shifts can render PU classifiers ineffective or even counterproductive, as they may learn patterns specific to historical data that no longer reflect current realities. This temporal degradation of model performance poses significant challenges for maintaining reliable classification systems.

Scalability and Computational Efficiency concerns become pronounced in modern applications involving high-dimensional data or large-scale datasets [39]. The risk decomposition inherent in PU learning requires multiple expectation calculations over different data subsets, increasing computational overhead compared to standard supervised learning [40]. For applications involving image, text, or sensor data with thousands of features, this additional computational burden can severely limit practical deployment, especially in resource-constrained environments or real-time applications where low latency and minimal memory footprint are critical requirements.

Table 2-2 Taxonomy of real-world challenges in PU learning applications

Challenge Category	Specific Manifestations	Impact on Learning	Domain Examples
Data Quality Issues	Label noise in positives; missing features; measurement errors	Biased risk estimation; overfitting to noise; reduced generalization	Medical diagnostics (misdiagnosis); social media (ambiguous content)
Distributional Problems	Covariate drift; concept drift; sample selection bias	Poor out-of-distribution performance; outdated decision boundaries	Financial fraud (evolving tactics); recommender systems (changing preferences)
Algorithmic Limitations	Sensitivity to hyperparameters; unstable optimization; slow convergence	Inconsistent results; high maintenance costs; difficult deployment	Industrial inspection (parameter tuning); autonomous systems (reliability concerns)
Theoretical-Practical Gaps	Violations of SCAR assumption; inaccurate prior estimation; model misspecification	Suboptimal performance; unreliable confidence estimates	Scientific discovery (novel patterns); security (adversarial environments)
Resource Constraints	Computational complexity; memory requirements; inference latency	Limited deployment scenarios; exclusion from edge computing	Mobile health applications; IoT sensor networks; real-time monitoring

(Source: Created by author)

Interpretability and Trustworthiness concerns have gained prominence as PU learning systems are increasingly deployed in high-stakes domains such as healthcare, finance, and criminal justice. The complex risk decomposition underlying PU methods creates opacity in decision-making processes, making it difficult to explain why a particular prediction was made or which features contributed most significantly [41]. This "black box" characteristic hinders adoption in regulated industries where model interpretability is not merely desirable but legally mandated.

As detailed in Table 2-2, these challenges are not merely theoretical concerns but manifest concretely across application domains, often interacting in complex ways that amplify their individual effects. The medical diagnostics domain illustrates this interplay clearly: label noise arises from diagnostic uncertainties and inter-rater variability; distributional shift occurs as diseases evolve and detection technologies advance; and interpretability requirements are stringent due to clinical accountability concerns. Similar patterns emerge in fraud detection, where adversaries actively adapt their strategies (creating distributional shift), ground truth is often ambiguous (creating label noise), and regulatory requirements demand explainable decisions.

The cumulative effect of these challenges is a significant gap between the theoretical promise of PU learning and its practical realization. While the field has made remarkable progress in formalizing learning frameworks and proving theoretical properties, the assumption

of clean, well-behaved data that underpins much of this work rarely holds in operational environments. This discrepancy highlights the urgent need for robust PU learning methods that maintain performance despite violated assumptions, a need that this research addresses directly through bounded loss function integration.

Current approaches to these challenges remain fragmented, with separate research threads addressing label noise, prior estimation, or distributional shift in isolation [42,43]. What is notably absent from the literature is a unified approach that enhances robustness across multiple dimensions through fundamental algorithmic design choices. The prevailing focus on developing increasingly complex estimators or correction mechanisms has overlooked the potential of simpler, more foundational improvements, particularly in loss function design. This gap represents both a limitation in current research and an opportunity for meaningful contribution, which this thesis seizes by demonstrating that bounded loss functions offer a principled, effective, and unifying strategy for enhancing PU learning robustness across diverse challenge scenarios.

2.2 Label Noise in Machine Learning

Label noise, the phenomenon where training data contains incorrect class annotations, represents one of the most pervasive and challenging problems in practical machine learning applications [44]. As machine learning systems have transitioned from controlled laboratory environments to real-world deployment, the assumption of clean, perfectly annotated data has proven increasingly untenable [45]. The ramifications of label noise extend across the entire machine learning pipeline, affecting data collection, model training, evaluation, and ultimately, deployment reliability. In essence, label noise creates a fundamental disconnect between the empirical distribution observed during training and the true underlying data distribution, leading models to learn and reinforce incorrect patterns.

The study of label noise has evolved from being considered a peripheral data quality issue to a central research topic in robust machine learning. Early approaches often treated label noise as an outlier detection problem or attempted to filter noisy examples through preprocessing [46]. However, contemporary understanding recognizes label noise as an inherent characteristic of many learning scenarios that must be addressed algorithmically throughout the training process [33]. This paradigm shift has given rise to an extensive literature on noise-robust learning, encompassing theoretical analyses of learnability under noise, algorithmic innovations for noise tolerance, and empirical evaluations across diverse application domains.

The particular significance of label noise in Positive-Unlabeled learning cannot be overstated. Unlike standard supervised learning where noise may affect both classes somewhat equally, PU learning exhibits asymmetric vulnerability because the entire learning signal depends crucially on the integrity of the (typically small) positive set [47]. A single mislabeled positive example can disproportionately influence risk estimation, especially when unbounded loss functions amplify its contribution [48]. This creates a compound challenge where the weaknesses of PU learning frameworks intersect with the general problem of label noise, producing particularly severe performance degradation.

2.2.1 Types and Impact of Label Noise

Label noise manifests in diverse forms, each with distinct characteristics, causes, and implications for learning algorithms. A comprehensive understanding requires categorization along multiple dimensions: the statistical relationship between true and observed labels, the dependence on instance features, and the contextual factors influencing noise generation.

Symmetric vs. Asymmetric Noise represents the most fundamental distinction in noise characterization. Symmetric or uniform noise occurs when each label has an equal probability of being flipped to any other class, independent of the true label [49]. This noise model, while mathematically tractable, often represents an oversimplification of real-world scenarios. More commonly, asymmetric or class-conditional noise occurs where label confusion follows predictable patterns, for instance, in medical diagnosis, borderline cases between similar diseases are more likely to be mislabeled than cases between distinctly different conditions. In PU learning, asymmetric noise is particularly relevant since mislabeling typically flows in one direction: true negatives incorrectly included in the positive set, while the converse (positives mislabeled as negative) is irrelevant since negatives are unlabeled by definition [50].

Instance-Dependent Noise introduces further complexity by linking noise probability to specific features of the data instance [51]. Unlike class-conditional noise where mislabeling probability depends only on the true class, instance-dependent noise recognizes that some examples are inherently more difficult to label accurately [52]. For example, in image classification, blurry or occluded images are more prone to labeling errors than clear, canonical examples. This form of noise is especially challenging because it creates a confounding relationship where the features most predictive of labeling difficulty may overlap with those predictive of class membership, making it difficult to disentangle noise patterns from genuine signal.

Contextual and Systematic Noise encompasses labeling errors that arise from specific contextual factors, such as subjective human biases, ambiguous labeling guidelines, or systematic errors in automated labeling pipelines [53]. In medical imaging, for instance, radiologist fatigue or variability in diagnostic criteria can introduce consistent patterns of mislabeling. In social media content moderation, evolving community standards and subjective interpretations create systematic labeling inconsistencies. These noise patterns often exhibit temporal or spatial correlations that violate the independent and identically distributed (i.i.d.) assumptions underlying most learning algorithms. The taxonomy of the label noise types and their properties are discussed in Table 2-3.

Table 2-3 Taxonomy of label noise types and their characteristics

Noise Type	Definition	Causal Mechanisms	Key Properties	Common Domains
Symmetric/Uniform	Each label has equal probability of being flipped to any other class	Random annotation errors; data entry mistakes	Stationary; class-balanced; analytically tractable	Large-scale crowdsourcing; automated data collection
Asymmetric/Class-Conditional	Label confusion follows class-specific patterns	Similarity between specific classes; systematic diagnostic confusion	Directional; class-imbalanced; reflects domain semantics	Medical diagnosis (similar diseases); product categorization (similar items)
Instance-Dependent	Noise probability depends on instance features	Inherently ambiguous examples; low-quality data samples	Feature-correlated; non-uniform; confounds with true signal	Image classification (blurry images); NLP (ambiguous phrasing)
Contextual/Systematic	Noise patterns arise from specific contextual factors	Human biases; ambiguous guidelines; systematic tool errors	Structured; often correlated; violates i.i.d. assumptions	Subjective ratings; evolving standards; multi-annotator settings
Adversarial Noise	Deliberate label corruption to mislead learning	Malicious data poisoning; strategic manipulation	Targeted; worst-case; security-critical	Security systems; spam filtering; adversarial environments

(Source: Created by author)

The impact of label noise on machine learning systems operates through multiple interconnected mechanisms, each contributing to the degradation of model performance and reliability. At the most fundamental level, label noise introduces bias in the empirical risk minimizer. When training data contains incorrect labels, the model that minimizes empirical risk on this corrupted data will generally differ from the model that would minimize risk on clean data.

Perhaps the most pernicious effect of label noise is its interaction with model capacity and complexity. Flexible, high-capacity models (such as deep neural networks) have sufficient expressive power to memorize noisy labels, leading to perfect training accuracy but poor generalization. This phenomenon, sometimes termed "label noise overfitting," is particularly problematic because it can create the illusion of successful learning while producing fundamentally unreliable models. The memorization of noisy labels is facilitated by the use of unbounded loss functions that assign disproportionately large penalties to incorrectly classified examples, causing the optimization process to prioritize fitting these outliers at the expense of learning robust patterns.

The statistical implications of label noise extend beyond simple bias to affect the entire learning dynamic. Noise increases the variance of parameter estimates, reduces effective sample size, and distorts confidence calibration [54]. In probabilistic models, label noise corrupts the likelihood function, leading to poorly calibrated posterior distributions [33]. In margin-based classifiers, noise reduces the effective margin and destabilizes the decision boundary [55]. These statistical effects compound in iterative learning processes, where initial errors can be amplified through feedback loops. Table 2-4 summarizes the multidimensional effect of label noise on machine learning systems.

Table 2-4 Multidimensional impact of label noise on machine learning systems

Impact Dimension	Specific Effects	Consequences for Learning	Particular Severity in PU Learning
Optimization Dynamics	Distorted gradients; unstable convergence; local minima attraction	Slower training; unreliable convergence; sensitivity to initialization	Amplified by unbounded loss functions; positive set noise dominates gradient
Statistical Properties	Biased risk estimates; increased variance; reduced effective samples	Poor generalization; overfitting to noise; unreliable confidence estimates	Compounded by risk decomposition; error propagation through prior estimation
Model Capacity Interaction	Overfitting to noise; memorization of corrupt patterns; reduced effective capacity	Perfect training accuracy with poor test performance; loss of robustness	High-capacity models may memorize all positives, including mislabeled ones
Decision Boundary Quality	Boundary distortion; margin reduction; increased ambiguity	Reduced discriminative power; higher error rates near boundary	Boundary shifts toward corrupted positives; unlabeled data provides weak correction
Downstream Applications	Poor transfer learning; unreliable confidence scores; security vulnerabilities	Limited deployment scenarios; increased maintenance costs	Compromises core value proposition of PU learning in noisy real-world settings

(Source: Created by author)

In the specific context of Positive-Unlabeled learning, the impact of label noise is magnified due to the framework's structural characteristics. The reliance on a small positive set means that even moderate noise rates can dramatically alter the empirical positive distribution. The risk estimation process, which involves subtracting positive expectations from marginal expectations, creates mathematical sensitivity to perturbations in the positive component. When the positive set contains mislabeled examples, the subtraction operation amplifies rather than cancels errors, potentially leading to severely distorted risk landscapes. Furthermore, the typical use of class prior information compounds the problem, if the class prior estimate is itself derived from noisy data, errors cascade through the entire learning process.

2.2.2 Noise-Robust Learning Methods

The growing recognition of label noise as a fundamental challenge has stimulated extensive research into noise-robust learning methods [56]. These approaches can be broadly categorized based on their underlying strategy: modifying the loss function to reduce sensitivity to noise, estimating and correcting the noise process, reweighting or filtering training examples, or designing architectures inherently resistant to label corruption. Each category represents a different philosophical approach to robustness, with varying trade-offs between theoretical guarantees, practical effectiveness, and implementation complexity.

Robust Loss Functions constitute perhaps the most direct approach to noise robustness. Rather than treating noise as a problem to be eliminated through preprocessing or modeling, these methods modify the learning objective itself to be inherently less sensitive to label errors. The theoretical foundation for this approach lies in the concept of noise-tolerant loss functions, those for which the minimizer under noisy data remains consistent with the minimizer under clean data under certain noise conditions. Early work established that the 0-1 loss possesses inherent noise robustness properties but is computationally intractable [57]. This insight spurred development of surrogate losses that approximate the noise tolerance of the 0-1 loss while remaining amenable to optimization.

The ramp loss represents a seminal contribution to robust loss design. By truncating the hinge loss beyond a certain margin, ramp loss limits the influence of any single example, preventing outliers from dominating the optimization. Formally defined as $\ell_{\text{ramp}}(y, f(x)) = \min(1, \max(0, 1 - yf(x)))$, this loss combines the convexity of hinge loss near the decision boundary with the boundedness of 0-1 loss for misclassified examples [58]. Variants like the savage loss and sigmoid ramp loss offer different trade-offs between differentiability and

robustness. More recently, symmetric losses have gained attention for their theoretical noise tolerance properties under symmetric noise conditions [59,60]. A loss function ℓ is symmetric if $\ell(y, f(x)) + \ell(-y, f(x)) = C$ for some constant C , ensuring that the expected loss under symmetric noise remains a linear function of the expected clean loss.

Generalized cross-entropy losses represent another important direction, attempting to combine the optimization-friendly properties of cross-entropy with the robustness of mean absolute error [61]. These losses typically introduce a shaping parameter that controls the trade-off between aggressiveness near the decision boundary and robustness to outliers. The truncated logistic loss used in this research belongs to this family, maintaining differentiability while bounding maximum loss contribution. Recent theoretical work has established formal conditions under which bounded losses provide robustness guarantees, particularly considering the link between loss Lipschitz continuity and noise tolerance [62].

Noise Modeling and Correction Methods take a fundamentally different approach by explicitly estimating the noise process and applying corrective transformations to the learning problem [63]. The noise transition matrix approach models the probability of observing label j given true label i , typically assuming this transition matrix is constant across instances. Once estimated (often using anchor points or clustering techniques), the transition matrix can be used to correct the loss function or modify the training labels. More sophisticated approaches model instance-dependent noise through meta-learning or two-network architectures where one network estimates noise patterns while another performs classification [64].

Data Reweighting and Selection Strategies attempt to identify and down-weight likely noisy examples during training. MentorNet train a separate network to assign weights to training examples based on estimated cleanliness [65]. Co-teaching methods maintain two networks that teach each other by selecting likely clean examples for updates. Curriculum learning approaches schedule training from easier to harder examples, often implicitly filtering noise in early stages [66]. These methods can be highly effective but introduce additional complexity and computational overhead.

Architectural Innovations for Robustness modify the model architecture itself to enhance noise resistance. Noise adaptation layers learn to model and correct label noise as part of the network. Confidence regularization techniques encourage the model to maintain low confidence on ambiguous or likely noisy examples. Label smoothing distributes probability

mass across classes rather than using one-hot encodings, reducing model overconfidence and sensitivity to exact label values [67]. Early stopping and regularization strategies, while not specifically designed for noise robustness, often provide practical benefits by preventing overfitting to noisy labels [68]. Table 2-5 compares the noise-robust learning paradigms.

Table 2-5 Comparative analysis of noise-robust learning paradigms

Paradigm	Core Mechanism	Theoretical Foundation	Practical Advantages	Limitations	Relevance to PU Learning
Robust Loss Functions	Modify loss to limit influence of noisy examples	Noise-tolerant loss theory; bounded loss properties	Simple implementation; minimal overhead; theoretical guarantees	May slow convergence; requires parameter tuning	Directly applicable; complements existing PU frameworks
Noise Modeling & Correction	Estimate and correct noise process	Statistical estimation theory; matrix completion	Can handle complex noise patterns; strong theoretical basis	Estimation errors propagate; requires anchor points	Challenging due to lack of negative labels for estimation
Data Reweighting & Selection	Identify and downweight noisy examples	Importance weighting; robust statistics	Adaptive to instance noise; can leverage curriculum learning	Computationally intensive; sensitive to reweighting heuristics	Limited by small positive set; may discard valuable positives
Architectural Modifications	Design models inherently resistant to noise	Representation learning; regularization theory	Integrated solution; can leverage deep learning advances	Architecture-specific; less generalizable	Requires rethinking PU-specific architectures
Ensemble & Committee Methods	Aggregate multiple models to reduce noise sensitivity	Variance reduction; committee machine theory	Reduces variance from noise; simple to implement	Computationally expensive; limited theoretical guarantees	Complements PU risk estimation; can ensemble multiple PU models

(Source: Created by author)

The application of these noise-robust methods to Positive-Unlabeled learning presents unique challenges and opportunities. Many standard noise robustness techniques rely on the availability of labeled examples from all classes for noise modeling or example weighting, a requirement fundamentally violated in PU settings. Methods based on noise transition matrices struggle because the absence of negative labels prevents estimation of the full transition matrix. Data reweighting approaches face difficulty because the small positive set offers limited redundancy for identifying mislabeled examples. These limitations highlight the need for

noise-robust methods specifically designed for the PU setting, or general methods that require minimal class-specific information.

Notably, robust loss functions emerge as particularly promising for PU learning applications due to their minimal requirements beyond the loss function itself. By modifying only the loss formulation, these methods can be directly integrated into existing PU frameworks like nnPU without requiring architectural changes, additional networks, or complex estimation procedures. The boundedness property of many robust losses aligns naturally with the stability requirements of PU risk estimation, potentially mitigating both the variance problems of uPU and the noise sensitivity of standard nnPU. This observation forms a key theoretical motivation for the approach developed in this thesis.

Notwithstanding extensive research on noise-robust learning, significant gaps remain in adapting these methods to weakly supervised settings generally and PU learning specifically. Most theoretical analyses assume fully labeled data, and most empirical evaluations use standard supervised benchmarks. The interaction between label noise and the distinctive risk estimation structure of PU learning remains underexplored theoretically, while practical implementations of noise-robust PU methods are scarce in the literature. This gap represents both a limitation in current research and an opportunity for meaningful contribution—one that this thesis addresses by developing and analyzing a bounded loss approach specifically designed for the unique challenges of noisy PU learning.

2.3 Bounded Loss Functions in Robust Learning

The design of loss functions represents a fundamental architectural choice in machine learning systems, governing not only what constitutes optimal performance but also how models respond to various forms of data imperfection. Among the diverse landscape of loss functions, bounded loss functions have emerged as particularly significant for robust learning due to their inherent capacity to limit the influence of individual data points, thereby providing natural resistance to outliers, mislabeled examples, and distributional anomalies. Unlike their unbounded counterparts, which can assign arbitrarily large penalties to incorrectly classified examples, bounded losses enforce a maximum contribution from any single instance, creating a form of "loss saturation" that prevents pathological examples from dominating the optimization landscape.

The theoretical foundation for bounded loss functions in robust learning rests upon several interconnected principles from statistical learning theory, optimization, and robust statistics. First, boundedness ensures that the loss function satisfies the conditions for uniform

convergence more readily, leading to tighter generalization bounds [69]. Second, by limiting the maximum gradient magnitude, bounded losses promote more stable optimization dynamics, particularly in stochastic gradient descent variants where extreme gradients can cause oscillation or divergence [70]. Third, from a robust statistics perspective, bounded losses correspond to M-estimators with redescending influence functions, meaning their influence on parameter estimates decreases for extreme residuals, exactly the property needed to mitigate the impact of label noise [71].

The mathematical formulation of boundedness is straightforward yet profound: a loss function ℓ is bounded if there exists a finite constant C such that $\ell(y, f(x)) \leq C$ for all possible inputs x , predictions $f(x)$, and true labels y . This simple constraint yields complex and beneficial behavioral consequences. When combined with appropriate regularization, bounded losses can prevent the memorization of noisy labels by making it "uneconomical" for the model to expend excessive capacity on fitting outliers. Instead, the optimization process naturally prioritizes patterns supported by multiple examples, effectively implementing a form of majority voting within the loss structure itself.

The development of bounded loss functions parallels the broader trend in machine learning toward robustness and reliability [72]. Early losses like the 0-1 loss were bounded but difficult to optimize. Later convex surrogates like hinge and logistic loss improved optimization but became unbounded. Modern research synthesizes these approaches, seeking losses that are both bounded for robustness and smooth for efficient optimization. This evolution reflects the growing understanding that models must be resilient under imperfect, real-world conditions, not just accurate under ideal assumptions.

2.3.1 Ramp Loss and Truncated Logistic Loss

Among the family of bounded loss functions, two specific formulations have demonstrated particular effectiveness in robust learning scenarios: the ramp loss and the truncated logistic loss. These losses exemplify different strategies for achieving boundedness while maintaining practical utility, offering complementary advantages that make them suitable for diverse applications.

The ramp loss, originally introduced in the context of robust support vector machines, represents perhaps the most conceptually straightforward approach to boundedness through direct truncation [73]. Mathematically defined as:

$$\ell_{\text{ramp}}(y, f(x)) = \min(1, \max(0, 1 - yf(x))) \quad (2-3)$$

This loss function can be understood as the hinge loss with its upper tail truncated at 1. The ramp loss exhibits three distinct behavioral regions: for correctly classified examples with sufficient margin ($yf(x) \geq 1$), the loss is zero; for examples near the decision boundary ($-1 < yf(x) < 1$), the loss decreases linearly with increasing margin; and for severely misclassified examples ($yf(x) \leq -1$), the loss saturates at a maximum value of 1. This saturation creates a crucial robustness property: no matter how far an example falls on the wrong side of the decision boundary, its contribution to the total loss cannot exceed a fixed ceiling.

The ramp loss possesses several theoretically desirable properties for robust learning. First, it satisfies the classification-calibrated condition, ensuring that minimizing the ramp risk leads to the Bayes optimal classifier. Second, its boundedness provides inherent resistance to outliers and adversarial examples, as these extreme instances cannot exert disproportionate influence. Third, the ramp loss exhibits the redescending influence property characteristic of robust M-estimators: as the residual $1 - yf(x)$ grows beyond the saturation point, the influence of the example on parameter estimates actually decreases. This counterintuitive property explains why ramp loss-based classifiers can effectively ignore blatantly mislabeled examples while focusing on ambiguous but potentially informative cases near the decision boundary.

Despite its non-convexity, the ramp loss admits efficient optimization through the Concave-Convex Procedure (CCCP) or specialized DC (Difference of Convex functions) programming algorithms. These methods decompose the ramp loss into a convex component (hinge loss) minus another convex component (a truncated hinge function), enabling iterative optimization with convergence guarantees. More recently, gradient-based methods with appropriate initialization have proven effective for optimizing ramp loss in deep learning contexts, particularly when combined with modern optimization techniques like Adam and learning rate scheduling.

The truncated logistic loss represents a different philosophical approach to boundedness, maintaining the smoothness and probabilistic interpretation of standard logistic loss while introducing an upper bound [74]. Formally defined as:

$$\ell_{\text{trunc-log}}(y, f(x)) = \min(\log(1 + e^{-yf(x)}), \tau) \quad (2-4)$$

where $\tau > 0$ is a truncation threshold, this loss behaves identically to standard logistic loss for moderately misclassified examples but caps the loss at τ for severely misclassified cases. The parameter τ thus controls the trade-off between robustness and sensitivity: lower τ values increase robustness by more aggressively limiting outlier influence, while higher τ values preserve more of the discriminative power of the standard logistic loss.

The truncated logistic loss as a whole offers several advantages over the ramp loss. Its smoothness ensures continuous gradients everywhere, facilitating optimization with standard first-order methods without requiring specialized decomposition algorithms [75]. The probabilistic interpretation inherent to logistic formulations remains largely intact, allowing for meaningful confidence estimates and straightforward integration with Bayesian methods. Perhaps most importantly, the truncated logistic loss provides a tunable robustness level through the τ parameter, enabling adaptation to different noise regimes or application requirements. This tunability makes it particularly suitable for scenarios where the noise characteristics may be unknown or variable.

Theoretical analysis reveals that the truncated logistic loss inherits many desirable properties from both its logistic and bounded lineages. It maintains the Fisher consistency of standard logistic loss for clean data while gaining robustness under contamination [76]. The boundedness ensures finite sample breakdown points that scale favorably with the truncation parameter. Moreover, the smooth transition at the truncation point prevents the sharp behavioral changes that can complicate optimization in piecewise losses like the ramp loss.

Table 2-6 Comparative analysis of bounded loss functions for robust learning

Property	Ramp Loss	Truncated Logistic Loss	Standard Logistic Loss	Hinge Loss
Boundedness	Yes (bounded between 0 and 1)	Yes (bounded by τ)	No (unbounded)	No (unbounded)
Convexity	Non-convex (difference of convex)	Non-convex due to truncation	Convex	Convex
Smoothness	Non-smooth at two points	Smooth everywhere	Smooth everywhere	Non-smooth at one point
Optimization Complexity	Requires CCCP or DC programming	Standard gradient methods	Standard gradient methods	Standard subgradient methods
Robustness to Label Noise	High (redescending influence)	Tunable via τ (high with small τ)	Low (susceptible to outliers)	Moderate (linear penalty)
Probabilistic Interpretation	No natural interpretation	Approximate probabilistic interpretation	Direct probabilistic interpretation	No natural interpretation
Theoretical Guarantees	Classification calibrated; finite breakdown point	Fisher consistent for clean data; robust with bounded τ	Fisher consistent; well-understood asymptotics	Margin maximization properties
Practical Tuning Parameters	None (fixed bound)	Truncation threshold τ	None (standard formulation)	None (standard formulation)

(Source: Created by author)

Beyond these canonical forms, recent research has developed numerous hybrid and adaptive bounded loss functions that attempt to combine the best properties of multiple approaches. The smooth ramp loss introduces smoothing around the kink points of the standard ramp loss to facilitate gradient-based optimization. The adaptive truncated loss dynamically adjusts the truncation threshold τ based on estimated noise levels or training progress. The generalized bounded exponential loss applies truncation to exponential loss functions, achieving robustness while maintaining the margin-enhancing properties of boosting algorithms. These innovations reflect an ongoing effort to balance the sometimes-competing demands of robustness, optimization efficiency, and theoretical tractability.

As summarized in Table 2-6, the choice between different bounded loss formulations involves trade-offs across multiple dimensions. The ramp loss offers maximum robustness with minimal tuning but requires specialized optimization. The truncated logistic loss provides smoother optimization and tunable robustness at the cost of an additional hyperparameter. Both approaches represent significant improvements over standard unbounded losses in noisy environments, yet neither has been systematically explored in the context of Positive-Unlabeled learning, a gap this research directly addresses.

2.3.2 Applications in Supervised and Weakly-Supervised Settings

The practical utility of bounded loss functions has been demonstrated across a diverse spectrum of machine learning paradigms, from traditional supervised classification to emerging weakly-supervised frameworks. These applications reveal both the versatility of bounded loss approaches and the domain-specific considerations that influence their effectiveness.

In fully supervised settings, bounded loss functions have proven particularly valuable in domains characterized by high label noise or frequent outliers. Medical imaging diagnostics represent a paradigmatic application area, where diagnostic uncertainty, inter-rater variability, and ambiguous clinical presentations introduce substantial label noise. Studies have shown that bounded loss functions can improve diagnostic accuracy by 8-15% compared to standard losses when label noise exceeds 20% [77]. Similarly, in financial fraud detection, where fraudulent patterns evolve rapidly and ground truth is often uncertain, bounded losses have demonstrated superior stability and reduced false positive rates.

The application of bounded losses to deep neural networks has generated particularly interesting findings. While early concerns suggested that the non-convexity of bounded losses might exacerbate optimization challenges in deep networks, empirical evidence has largely

contradicted this assumption. Properly implemented bounded losses often improve training stability in deep networks by preventing gradient explosion from mislabeled examples. Moreover, the implicit regularization effect of boundedness can complement explicit regularization techniques like dropout and weight decay, sometimes reducing the need for extensive hyperparameter tuning. In large-scale image classification benchmarks with simulated label noise, bounded loss variants have consistently outperformed standard cross-entropy, with improvements most pronounced at higher noise levels (30-50% corruption) [78].

Bounded loss functions have shown promising synergies with semi-supervised learning (SSL) methods, particularly those based on consistency regularization. In SSL frameworks like MixMatch and FixMatch, bounded losses applied to labeled examples can prevent noise propagation to the unlabeled pool through pseudo-labeling mechanisms [79]. The boundedness ensures that severely mislabeled examples do not generate confidently incorrect pseudo-labels that would then contaminate training on unlabeled data. This "noise containment" property is especially valuable in real-world SSL scenarios where labeled data is both scarce and noisy, precisely the conditions where SSL is most needed yet most vulnerable.

Recent research has developed hybrid loss formulations that combine bounded losses for labeled data with standard consistency losses for unlabeled data [80,81]. These approaches recognize that different loss characteristics may be appropriate for different data subsets: boundedness provides robustness for potentially noisy labeled examples, while unbounded consistency terms can encourage aggressive exploration of the unlabeled manifold. Theoretical analysis suggests that such hybrid approaches can achieve better trade-offs between robustness and utilization of unlabeled data than either purely bounded or purely standard approaches.

Beyond PU learning, bounded loss functions have found applications in various other weakly-supervised settings. In multiple-instance learning (MIL), where labels are assigned to bags of instances rather than individual examples, bounded losses help mitigate the effect of mislabeled bags and improve instance-level classification. In partial label learning, where each example is associated with a set of candidate labels containing the true label, bounded losses reduce the penalty for incorrect candidates, preventing overemphasis on disambiguation at the expense of discriminative learning.

Perhaps most relevantly to this research, bounded losses have shown promise in noisy label learning with fully supervised but corrupted data. Here, the theoretical connection between loss boundedness and noise tolerance has been most rigorously established. Under symmetric label noise, any loss function satisfying certain symmetry conditions (including

appropriately bounded variants) provably maintains consistency, the minimizer on noisy data equals the minimizer on clean data. Even under asymmetric noise, bounded losses with appropriate truncation levels can significantly reduce error compared to unbounded alternatives.

Table 2-7 Application domains and performance characteristics of bounded loss functions

Application Domain	Primary Data Challenges	Bounded Loss Benefits	Empirical Performance Gains	Key Implementation Considerations
Medical Diagnosis	Label noise from diagnostic uncertainty; class imbalance	Reduces overfitting to misdiagnosed cases; improves specificity	8-15% accuracy improvement at >20% noise	Integration with domain-specific architectures; calibration requirements
Financial Fraud Detection	Evolving fraud patterns; ambiguous ground truth	Stabilizes learning against novel attack patterns; reduces false positives	20-30% reduction in false positives at similar recall	Real-time processing constraints; concept drift adaptation
Autonomous Driving	Ambiguous scenarios; safety-critical errors	Limits influence of corner cases; improves worst-case performance	15-25% improvement on rare but critical scenarios	Certification requirements; interpretability needs
Social Media Content Moderation	Subjective labeling; evolving community standards	Accommodates labeling inconsistencies; adapts to policy changes	10-20% better consistency across annotator groups	Scalability to massive datasets; multilingual adaptation
Scientific Discovery	Noisy experimental measurements; hypothesis refinement	Robust to measurement errors; prevents overinterpretation	Improved reproducibility in downstream analysis	Integration with domain knowledge; small sample adaptation
Manufacturing Quality Control	Visual inspection ambiguities; defect variability	Tolerates borderline cases; reduces false rejections	25-35% reduction in false rejection rates	Integration with production lines; hardware constraints

(Source: Created by author)

The theoretical analysis of bounded loss functions in weakly-supervised settings remains an active research area with important open questions. While the noise tolerance properties of bounded losses are well-established for symmetric noise in supervised settings, their behavior under the asymmetric noise patterns typical of PU learning (where only false positives occur in the labeled set) requires more nuanced analysis. Recent work has begun to characterize the conditions under which bounded losses maintain consistency in such asymmetric scenarios, with preliminary results suggesting that appropriate bounding can indeed mitigate bias from one-sided noise.

As summarized in Table 2-7, the empirical success of bounded loss functions across diverse application domains underscores their practical value while highlighting domain-

specific implementation considerations. What remains notably absent from this landscape is a systematic exploration of bounded losses in the specific context of Positive-Unlabeled learning, notwithstanding the particularly severe consequences of label noise in PU settings and the natural alignment between boundedness and PU's stability requirements. This research gap is especially striking given the parallel developments in bounded loss theory and PU learning methodology over the past decade. The present work addresses this gap by not only applying bounded losses to PU learning but also developing a comprehensive theoretical framework that characterizes their behavior under the unique statistical conditions of PU learning with label noise.

An important theoretical limitation concerns the interaction between boundedness and model capacity. While bounded losses prevent overfitting to noise in finite-capacity models, extremely high-capacity models (like modern deep networks) may still memorize bounded losses through parameter redundancy. This has led to the development of capacity-aware bounding strategies that adjust loss bounds based on estimated model complexity or training progress. Another theoretical challenge involves the calibration properties of bounded loss minimizers, while they may achieve good classification accuracy, their probability estimates may be poorly calibrated, requiring post-hoc calibration techniques for applications requiring reliable confidence scores [82].

The integration of bounded loss functions into PU learning represents more than a simple technical combination; it embodies a principled approach to addressing one of the most fundamental challenges in weakly supervised learning. By limiting the influence of any single potentially mislabeled positive example, bounded losses provide a natural mechanism for enhancing the robustness of PU risk estimators without requiring complex noise modeling or additional supervision. This approach leverages insights from multiple branches of machine learning, robust statistics, optimization theory, and statistical learning theory to create a solution that is both theoretically grounded and practically effective. The following chapters develop this integration in detail, providing formal analysis, algorithmic implementation, and empirical validation of bounded loss PU learning across diverse noise conditions and application scenarios.

2.4 Research Gaps in Noise-Robust PU Learning

The comprehensive review of existing literature reveals significant advances across the interrelated domains of Positive-Unlabeled learning, label noise robustness, and bounded loss function design. However, a critical synthesis of these developments exposes several important

gaps at their intersection, gaps that limit the practical applicability of PU learning in real-world scenarios where data imperfections are inevitable. These research gaps represent not merely incremental opportunities for improvement but fundamental limitations in current approaches to noise-robust weakly supervised learning.

2.4.1 Theoretical Gap: Absence of Formal Robustness Guarantees for PU Learning Under Label Noise

Not overlooking the extensive theoretical development in PU learning frameworks, particularly regarding risk consistency, estimation bias, and convergence properties, the theoretical literature remains strikingly silent on the formal analysis of robustness under label contamination. The foundational SCAR assumption, while mathematically convenient, creates a brittle theoretical foundation that collapses when even moderate label noise is introduced. Current theoretical analyses of PU learning typically proceed under the idealized assumption of perfectly labeled positives, leaving unanswered critical questions about the behavior of PU risk estimators when this assumption is violated.

1) **Mathematical Incompleteness in Risk Analysis:** Existing theoretical guarantees for PU learning including the unbiasedness of uPU and the consistency of nnPU are established under noise-free conditions. When label noise is introduced, the risk decomposition in Equation (3) becomes mathematically invalid, as the expectation over the observed positive distribution no longer equals the expectation over the true positive distribution. The propagation of this error through the risk estimator remains uncharacterized theoretically, creating uncertainty about whether small amounts of label noise led to bounded performance degradation or catastrophic failure.

2) **Missing Characterization of Noise-Induced Bias-Variance Tradeoffs:** In supervised learning with label noise, theoretical analyses often characterize the bias-variance decomposition of the risk estimator under contamination. For PU learning, no analogous characterization exists. It remains unknown how label noise affects the delicate balance between the bias introduced by the non-negative rectification in nnPU and the variance inherent in the difference of expectations. This theoretical gap prevents principled algorithm design and parameter selection for noisy PU scenarios.

2.4.2 Methodological Gap

The methodological landscape of PU learning reveals a striking dichotomy: sophisticated frameworks for handling unlabeled data co-exist with primitive approaches to label noise.

While PU learning has matured considerably in its ability to leverage unlabeled information, it has largely failed to incorporate advances from the parallel field of noise-robust learning. This methodological segregation represents a critical gap that prevents the deployment of PU learning in realistic settings.

1) **Inadequate Transfer of Noise-Robust Techniques:** The rich literature on noise-robust learning encompassing robust loss functions, noise modeling, data reweighting, and architectural innovations has seen minimal integration into PU learning frameworks. Most PU methods continue to employ standard unbounded loss functions (logistic, hinge, exponential) that are demonstrably sensitive to label noise. The few exceptions that consider label noise typically adopt ad hoc solutions rather than principled integration of established robust learning principles.

2) **Overreliance on Two-Stage Approaches:** The predominant strategy for handling label noise in PU learning involves sequential processing: first identify and filter potentially noisy positives, then apply standard PU learning to the presumably cleaned set. This approach suffers from multiple limitations: error propagation from the filtering stage, inefficient use of data, and failure to account for the continuum of noise severity. More fundamentally, it represents an avoidance of the core problem rather than a solution to it.

3) **Neglect of Loss Function Design as a Robustness Mechanism:** Perhaps the most significant methodological gap lies in the nearly universal disregard for loss function design as a mechanism for achieving noise robustness in PU learning. While bounded loss functions have demonstrated remarkable effectiveness in supervised noisy label learning, their potential application to PU settings remains almost entirely unexplored. This omission is particularly striking given the natural alignment between the boundedness property (which limits individual example influence) and the stability requirements of PU risk estimation (which involves delicate differences of expectations).

2.4.3 Empirical Gap: Limited Validation Under Realistic Noise Conditions

The empirical evaluation of PU learning methods reveals systematic biases in experimental design that limit their relevance to real-world applications. These biases create an empirical gap between laboratory validation and practical deployment, obscuring the true limitations of existing methods and hindering progress toward robust solutions.

1) **Synthetic Noise Models vs. Real-World Noise Patterns:** Most empirical studies of PU learning employ artificially generated symmetric label noise, typically implemented by randomly flipping positive labels to negative. While mathematically convenient, this model

fails to capture the structured, asymmetric, and instance-dependent nature of real-world label noise. In medical diagnosis, for instance, mislabeling follows specific confusion patterns between similar conditions; in fraud detection, labeling errors correlate with transaction complexity and novelty. The absence of standardized benchmarks incorporating realistic noise patterns prevents meaningful comparison of robustness claims.

2) **Narrow Performance Evaluation:** Empirical evaluations of PU learning overwhelmingly focus on aggregate metrics like accuracy and AUC, neglecting crucial dimensions of robustness. There is minimal reporting of performance breakdowns across noise levels, sensitivity analyses to noise rate misspecification, or characterization of failure modes under increasing contamination. This narrow evaluation creates an incomplete picture of method suitability for different operational environments.

3) **Limited Domain Coverage:** The empirical validation of PU learning remains concentrated in a few domains (primarily text classification and bioinformatics), with limited exploration in areas where label noise is particularly problematic, such as medical imaging, autonomous systems, or industrial quality control. This domain concentration creates uncertainty about method generalizability and obscures domain-specific challenges that might necessitate tailored solutions. Table 2-8 outlines the recent empirical limitations in PU learning

Table 2-8 Empirical limitations in current PU learning research

Evaluation Dimension	Current Practice	Ideal Practice	Consequences of Gap
Noise Modeling	Symmetric random flipping	Realistic asymmetric and instance-dependent noise	Overestimation of method robustness
Performance Metrics	Aggregate accuracy/AUC	Noise-level-stratified metrics, stability measures	Incomplete understanding of operational envelopes
Dataset Diversity	Limited domains (text, bioinformatics)	Cross-domain validation (vision, audio, tabular)	Uncertain generalizability to new applications
Baseline Comparison	Comparison only to other PU methods	Comparison to noise-robust supervised methods	Inability to assess opportunity cost of PU approach
Hyperparameter Reporting	Incomplete or inconsistent	Full sensitivity analysis and ablation studies	Difficulty in reproduction and fair comparison

(Source: Created by author)

2.4.4 Integrative Gap: Disconnected Research Trajectories

Perhaps the most fundamental gap in current research lies in the disconnection between three robust research trajectories: PU learning methodology, noise-robust learning theory, and

bounded loss function design. These domains have evolved largely independently, with minimal cross-pollination of ideas, despite their natural complementarity and shared objectives.

1) Methodological Isolationism: Researchers in PU learning have primarily focused on improving risk estimation and leveraging unlabeled data, often treating label quality as an external concern. Conversely, researchers in noise-robust learning have concentrated on supervised settings, with limited consideration for the unique challenges of weakly supervised scenarios. This disciplinary siloing has prevented the synthesis of insights that could lead to more fundamentally robust learning paradigms.

2) Theoretical Fragmentation: The theoretical foundations of PU learning (statistical risk estimation), noise-robust learning (influence functions, breakdown points), and bounded loss design (classification calibration, Fisher consistency) employ different mathematical frameworks and assumptions. This fragmentation has hindered the development of unified theoretical perspectives that could provide holistic guidance for algorithm design under simultaneous challenges of missing labels and label noise.

3) Lack of Unified Design Principles: The absence of integrative research has resulted in a methodological landscape characterized by ad hoc solutions rather than principled approaches. There exist no established design principles for creating PU learning methods that are inherently robust to label noise, nor systematic frameworks for adapting noise-robust techniques to the PU setting. This represents a significant missed opportunity for creating more reliable weakly supervised learning systems.

2.4.5 Bridging the Gaps: The Contribution of This Research

The research presented in this thesis directly addresses these interconnected gaps through a unified approach that integrates bounded loss function design into the PU learning paradigm. By systematically incorporating ramp loss and truncated logistic loss into the nnPU framework, this work bridges theoretical, methodological, and empirical divides to create a more robust approach to learning from positive and unlabeled data.

1) Theoretical Integration: This research develops a comprehensive theoretical framework that characterizes the robustness properties of bounded loss functions in PU learning. By establishing formal bounds on the influence of label noise and proving consistency properties under bounded losses, it addresses the theoretical gap in noise-robust PU learning theory.

2) Methodological Synthesis: The proposed framework represents a principled synthesis of three research traditions: PU risk estimation from the PU learning literature, bounded loss

design from robust statistics, and noise tolerance analysis from robust machine learning. This synthesis creates a method that is both theoretically grounded and practically implementable.

3) Empirical Advancement: Through extensive experimentation across synthetic and real-world datasets with varied noise conditions, this research addresses the empirical gaps in current PU learning evaluation. The experiments systematically explore performance across noise levels, compare against appropriate baselines, and validate on diverse domains including medical diagnosis and image classification.

4) Integrative Contribution: Most fundamentally, this research demonstrates the value of integrative thinking across traditionally separate research domains. By showing that bounded loss functions, a concept from robust supervised learning can effectively address label noise challenges in PU learning, a weakly supervised paradigm which establishes a template for future cross-disciplinary innovation in robust machine learning.

The identified research gaps collectively represent not merely academic oversights but practical impediments to the deployment of PU learning in real-world applications where data imperfections are the norm. By systematically addressing these gaps, this research advances the state of the art toward more reliable, robust, and practical weakly supervised learning systems capable of operating effectively in the imperfect data environments that characterize most real-world applications.

2.5 Chapter Summary

This chapter provides a comprehensive examination of the theoretical foundations and methodological developments in Positive-Unlabeled (PU) learning, label noise robustness, and bounded loss functions, culminating in the identification of critical research gaps that motivate this thesis.

Section 2.1 introduces PU learning as a weakly supervised paradigm where only positive examples are labeled while the remaining data lacks class identifiers. The field evolved from text classification to formal mathematical frameworks operating under the Selected Completely at Random (SCAR) assumption. Two primary methodological approaches exist: risk estimation and two-step strategies. The Unbiased PU (uPU) framework offers theoretical elegance through unbiased risk estimation but suffers from high variance and negative empirical risks. The Non-negative PU (nnPU) framework addresses these limitations through a rectification operator, becoming the standard for modern applications. Real-world challenges include data quality

issues (label noise), distributional problems (covariate drift), algorithmic limitations, theoretical-practical gaps (SCAR violations), and resource constraints.

Section 2.2 examines label noise as a pervasive challenge that creates disconnects between empirical and true data distributions. A taxonomy distinguishes symmetric versus asymmetric noise, instance-dependent versus class-conditional patterns, and contextual versus adversarial corruption. Noise impacts learning through biased risk estimation, increased variance, reduced effective sample size, and distorted decision boundaries. In PU learning, these effects amplify because risk decomposition involves subtracting positive expectations from marginal expectations, creating mathematical sensitivity to positive component perturbations.

Section 2.3 explores bounded loss functions as principled robust learning tools. Unlike unbounded losses, bounded variants enforce maximum contribution from any single instance, preventing pathological examples from dominating optimization. The ramp loss truncates hinge loss at 1, exhibiting three regions: zero loss for correct classifications, linear decrease near the decision boundary, and saturation for severe misclassifications. Its redescending influence property means parameter impact decreases for extreme residuals. The truncated logistic loss maintains smoothness and probabilistic interpretation while introducing an upper bound, offering tunable robustness via threshold τ . Applications across medical diagnosis, fraud detection, and autonomous driving demonstrate 8-35% performance improvements in noisy environments.

Section 2.4 synthesizes literature streams to identify critical research gaps. First, no systematic integration of bounded loss functions into PU learning frameworks exists despite natural alignment with stability requirements. Second, theoretical behavior of bounded losses under asymmetric noise patterns characteristic of PU learning remains uncharacterized. Third, empirical validation of noise-robust PU methods across diverse conditions is scarce. Fourth, practical guidelines for selecting bounded loss parameters in PU contexts are absent. Fifth, interaction between loss boundedness and class prior estimation uncertainty remains unexplored. These gaps motivate this thesis's central contribution: developing and analyzing a bounded loss approach specifically designed for noisy PU learning, bridging theoretical foundations with practical effectiveness.

Chapter III METHODOLOGY

3.1 Introduction

This chapter presents the methodological foundation of our proposed robust Positive-Unlabeled learning framework. Building upon the identified research gaps discussed in Chapter 2, we develop a comprehensive approach that integrates bounded loss functions into the PU learning paradigm to address the critical challenge of label noise in the positive set. The methodology presented here is designed to achieve three interconnected objectives: theoretical robustness, algorithmic efficiency, and practical applicability.

Our approach is grounded in the recognition that label noise in PU learning represents not merely a nuisance to be mitigated through preprocessing, but a fundamental challenge requiring integrated solutions at the algorithmic level. Rather than treating noise as an external corruption to be filtered or corrected, we reconceptualize it as an intrinsic property of the learning problem that must be addressed through appropriate loss function design and risk estimation. This perspective leads us to develop a framework that maintains the theoretical advantages of established PU methods while enhancing their resilience to label contamination.

The methodological development proceeds through several logically connected stages. First, we formalize the PU learning problem and its standard risk estimation frameworks, establishing notation and assumptions that will be used throughout. Second, we introduce a realistic label noise model that captures the contamination of positive labels, a scenario commonly encountered in practical applications. Third, we develop our core innovation: the integration of bounded loss functions into the PU risk estimation framework. Fourth, we provide theoretical analysis establishing the robustness properties of our approach. Finally, we present an algorithmic instantiation suitable for practical implementation.

The methodological choices made in this chapter are guided by principles of theoretical soundness, ensuring that our extensions maintain the desirable properties of existing PU frameworks; practical feasibility, ensuring that the resulting algorithms are computationally tractable and implementable with standard optimization techniques; and robustness by design, embedding noise tolerance directly into the learning objective rather than treating it as an external consideration. By adhering to these principles, we aim to develop a framework that not only addresses the specific challenge of label noise but also advances the broader goal of creating more reliable weakly supervised learning systems.

3.2 Preliminaries and Problem Formulation

3.2.1 PU Learning Setup and Risk Estimation

We begin by formalizing the standard Positive-Unlabeled learning problem within a probabilistic framework. Let $\mathcal{X} \subseteq \mathbb{R}^d$ denote the input feature space and $\mathcal{Y} = \{+1, -1\}$ the binary label space, where $+1$ indicates the positive class and -1 the negative class. In PU learning, we observe two distinct datasets:

1) Labeled positive examples: $\mathcal{P} = \{x_i^p\}_{i=1}^{n_p}$ independently and identically distributed (i.i.d.) according to the true positive conditional distribution $p(x | Y = +1)$.

2) Unlabeled examples: $\mathcal{U} = \{x_i^u\}_{i=1}^{n_u}$ i.i.d. according to the marginal distribution $p(x)$, which represents a mixture of positive and negative instances.

The marginal distribution can be expressed as:

$$p(x) = \pi p(x | Y = +1) + (1 - \pi) p(x | Y = -1) \quad (3-1)$$

where $\pi = P(Y = +1)$ denotes the class prior, representing the proportion of positive instances in the overall population. This prior is typically assumed to be known or estimable from the data.

The learning objective is to train a binary classifier $f: \mathcal{X} \rightarrow \mathbb{R}$ that predicts the label of a new instance by minimizing the expected risk:

$$R(f) = \mathbb{E}_{(x,y) \sim p(x,y)} [\ell(y, f(x))] \quad (3-2)$$

where $\ell: \mathcal{Y} \times \mathbb{R} \rightarrow \mathbb{R}^+$ is a surrogate loss function (e.g., logistic loss, hinge loss) that approximates the 0-1 loss.

Since negative labels are unavailable in the PU setting, the true risk cannot be computed directly. The key innovation of modern PU learning lies in reformulating this risk using only observable quantities. The unbiased PU (uPU) risk estimator achieves this through algebraic manipulation:

$$\begin{aligned} R_{uPU}(f) &= \pi \cdot \mathbb{E}_{x \sim p(x|Y=+1)} [\ell(+1, f(x))] \\ &\quad + \mathbb{E}_{x \sim p(x)} [\ell(-1, f(x))] \\ &\quad - \pi \cdot \mathbb{E}_{x \sim p(x|Y=+1)} [\ell(-1, f(x))] \end{aligned} \quad (3-3)$$

This estimator is unbiased with respect to the true risk under the Selected Completely at Random (SCAR) assumption, meaning that $\mathbb{E}[R_{uPU}(f)] = R(f)$. However, it suffers from

practical limitations, most notably the potential for negative empirical risk estimates and high variance, particularly when the classifier is flexible or the loss function is unbounded.

To address these issues, the non-negative PU (nnPU) risk estimator introduces a rectification mechanism:

$$\begin{aligned} R_{nnPU}(f) &= \pi \cdot \mathbb{E}_{x \sim p(x|Y=+1)}[\ell(+1, f(x))] \\ &\quad + \max(0, \mathbb{E}_{x \sim p(x)}[\ell(-1, f(x))]) \\ &\quad - \pi \cdot \mathbb{E}_{x \sim p(x|Y=+1)}[\ell(-1, f(x))] \end{aligned} \quad (3-4)$$

The $\max(0, \cdot)$ operator ensures non-negative empirical risk estimates, stabilizing training while maintaining consistency properties. This formulation forms the foundation upon which we build our robust extension.

Both uPU and nnPU typically employ standard unbounded loss functions such as logistic loss $\ell_{\log}(y, f(x)) = \log(1 + \exp(-yf(x)))$ or hinge loss $\ell_{\text{hinge}}(y, f(x)) = \max(0, 1 - yf(x))$. As discussed in Chapter 2, these unbounded losses are particularly sensitive to label noise, as mislabeled examples can exert disproportionately large influence on the risk estimate and its gradient.

3.2.2 Label Noise Model

To formally address label noise in PU learning, we must first establish a mathematical model that captures the contamination process. We focus on noise in the positive set, as this represents the most critical vulnerability in PU learning, negative examples are by definition unlabeled and thus cannot be mislabeled in the traditional sense.

Let $\tilde{p}(x | Y = +1)$ denote the observed noisy positive distribution, which may differ from the true positive distribution $p(x | Y = +1)$ due to labeling errors. We model this corruption using a mixture model:

$$\tilde{p}(x | Y = +1) = (1 - \eta) \cdot p(x | Y = +1) + \eta \cdot p(x | Y = -1) \quad (3-5)$$

where $\eta \in [0, 1]$ represents the noise rate, the proportion of true negative instances incorrectly included in the positive set. This model assumes that with probability $1 - \eta$, a true positive is correctly labeled, while with probability η , a true negative is mislabeled as positive.

Several important properties characterize this noise model:

1) Unidirectional contamination: Noise flows only from the negative class to the positive class, reflecting the practical reality that negative instances are unavailable for direct labeling in PU settings.

2) Uniformity assumption: The noise rate η is constant across the feature space, meaning all negative instances have equal probability of being mislabeled. While this assumption simplifies analysis, we note that real-world noise may exhibit instance-dependent characteristics, an extension we discuss in Chapter 5.

3) Preservation of marginal distribution: The overall marginal distribution $p(x)$ remains unchanged despite label noise, as mislabeling merely reallocates samples between the observed positive set and the unlabeled set.

Under this noise model, the empirical positive set $\mathcal{P}$ is actually drawn from $\tilde{p}(x | Y = +1)$ rather than $p(x | Y = +1)$. Consequently, standard PU risk estimators become biased. For the nnPU estimator (Equation 3.4), the observed risk becomes:

$$\begin{aligned} \tilde{R}_{nnPU}(f) &= \pi \cdot \mathbb{E}_{x \sim \tilde{p}(x|Y=+1)}[\ell(+1, f(x))] \\ &\quad + \max(0, \mathbb{E}_{x \sim p(x)}[\ell(-1, f(x))] \\ &\quad - \pi \cdot \mathbb{E}_{x \sim \tilde{p}(x|Y=+1)}[\ell(-1, f(x))]) \end{aligned} \quad (3-6)$$

The difference between the noisy risk $\tilde{R}_{nnPU}(f)$ and the true risk $R_{nnPU}(f)$ quantifies the impact of label noise on learning. This difference can be decomposed into two components corresponding to the two expectations involving the positive distribution:

$$\begin{aligned} \Delta R(f) &= \tilde{R}_{nnPU}(f) - R_{nnPU}(f) \\ \Delta R(f) &= \pi \cdot \mathbb{E}_{x \sim \Delta p(x)}[\ell(+1, f(x)) - \ell(-1, f(x))] \\ &\quad + \text{non-linear effects from max operator} \end{aligned} \quad (3-7)$$

where $\Delta p(x) = \tilde{p}(x | Y = +1) - p(x | Y = +1) = \eta \cdot [p(x | Y = -1) - p(x | Y = +1)]$.

The magnitude of this risk distortion depends critically on the loss function properties. For unbounded losses like logistic loss, $\Delta R(f)$ can be arbitrarily large, as mislabeled negatives far from the decision boundary can produce extreme loss values. For bounded losses, however, the impact is inherently limited, providing the theoretical motivation for our approach.

We further define the signal-to-noise ratio (SNR) in the positive set as:

$$\text{SNR} = \frac{(1 - \eta) \cdot \mathbb{E}_{x \sim p(x|Y=+1)}[\|x - \mu_+\|^2]}{\eta \cdot \mathbb{E}_{x \sim p(x|Y=-1)}[\|x - \mu_-\|^2]} \quad (3-8)$$

where μ_+ and μ_- are the class means. This ratio quantifies the relative strength of the true positive signal compared to the noise introduced by mislabeled negatives. Learning becomes increasingly challenging as SNR decreases, necessitating more robust approaches.

The noise model presented here provides a formal framework for analyzing and mitigating the effects of label contamination in PU learning. In the following sections, we develop our robust approach based on the insight that bounded loss functions naturally limit the risk distortion $\Delta R(f)$, thereby enhancing classifier stability and performance under label noise.

3.3 Proposed Framework: Bounded Loss PU Learning

3.3.1 Motivation for Bounded Loss Functions

The vulnerability of traditional PU learning frameworks to label noise stems fundamentally from their reliance on unbounded loss functions. Standard surrogate losses such as logistic, hinge, and exponential losses share the property that as the margin $yf(x) \rightarrow -\infty$ (i.e., as an example is severely misclassified), the loss value grows without bound: $\ell(y, f(x)) \rightarrow \infty$. This characteristic, while beneficial for encouraging large-margin classification in clean settings, becomes detrimental when labels are corrupted.

In the presence of label noise, unbounded losses create a pathological learning dynamic. Consider a negative instance x^- that has been mislabeled as positive. During training, the classifier will naturally tend to assign a negative score $f(x^-) < 0$ to this example. Since it carries a positive label $y = +1$, the product $yf(x^-)$ becomes negative, potentially with large magnitude. An unbounded loss function responds to this scenario by assigning an excessively large penalty, causing the optimization process to prioritize correcting this single mislabeled example over learning patterns from the majority of correctly labeled data.

The mathematical consequence is an inflated and distorted risk landscape. The empirical risk becomes dominated by contributions from mislabeled outliers, particularly those far from the decision boundary. This distortion manifests in two ways:

- 1) Gradient explosion: The gradient $\nabla_{\theta} \ell(y, f_{\theta}(x))$ for mislabeled examples can become disproportionately large, destabilizing optimization and causing erratic parameter updates.
- 2) Risk bias: The empirical risk estimate deviates significantly from the true population risk, with the deviation increasing with both noise rate and loss unboundedness.

Bounded loss functions address this fundamental limitation by enforcing an upper limit on the contribution of any single example. Formally, a loss function ℓ_b is bounded if there exists a finite constant C such that:

$$\sup_{(x,y) \in \mathcal{X} \times \mathcal{Y}} \ell_b(y, f(x)) \leq C < \infty \forall f \in \mathcal{F} \quad (3-9)$$

This boundedness property ensures that even severely mislabeled examples cannot exert excessive influence on the overall risk estimate or its gradient. In the PU learning context specifically, bounded losses provide three key advantages:

1) Natural Noise Robustness: By capping the maximum penalty for any example, bounded losses prevent mislabeled positives from dominating the risk estimation. This is particularly valuable in PU learning where the positive set is typically small and thus more vulnerable to contamination effects.

2) Stabilized Risk Decomposition: The nnPU risk estimator involves subtracting two potentially large expectations. See Eq. (3-4). When unbounded losses cause these expectations to be large, their difference becomes numerically unstable. Bounded losses constrain both terms, making the subtraction more stable and reliable.

3) Improved Gradient Dynamics: The boundedness of the loss implies bounded gradients, leading to smoother optimization trajectories and reduced sensitivity to learning rate choices. This is especially beneficial in stochastic optimization where extreme gradients can cause oscillation or divergence. Table 3-1 compares the loss function features for noisy PU learning.

Table 3-1 Comparison of loss function properties for noisy PU learning

Property	Unbounded Losses (e.g., Logistic)	Bounded Losses (e.g., Ramp)	Implication for Noisy PU Learning
Maximum Penalty	Unbounded ($\rightarrow\infty$)	Bounded ($\leq C$)	Prevents outlier domination in risk estimation
Gradient Magnitude	Can be arbitrarily large	Bounded by Lipschitz constant	Stabilizes optimization; reduces variance
Risk Decomposition Stability	Prone to cancellation errors	Numerically stable differences	More reliable nnPU risk computation
Noise Sensitivity	High (amplifies mislabeling effects)	Low (caps mislabeling impact)	Maintains performance under label contamination
Theoretical Guarantees	Consistency under clean labels	Consistency + robustness bounds	Provable performance under noise

(Source: Created by author)

The motivation for bounded losses in PU learning is thus both theoretical and practical. Theoretically, boundedness provides inherent robustness properties that can be formally characterized. Practically, it leads to more stable training dynamics and better performance in real-world scenarios where label noise is unavoidable. Our framework leverages these advantages while maintaining compatibility with established PU learning methodology.

3.3.2 Ramp Loss and Truncated Logistic Loss Formulations

We now define two specific bounded loss functions that form the core of our robust PU learning framework: the ramp loss and the truncated logistic loss. These losses represent complementary approaches to achieving boundedness while maintaining desirable optimization properties.

The ramp loss is a piecewise linear bounded surrogate that combines the hinge loss's margin-maximizing property with the 0-1 loss's boundedness. It is defined as:

$$\ell_{\text{ramp}}(y, f(x)) = \begin{cases} 1 & \text{if } yf(x) \leq -1 \\ 1 - yf(x) & \text{if } -1 < yf(x) < 1 \\ 0 & \text{if } yf(x) \geq 1 \end{cases} \quad (3-10)$$

This formulation can be more compactly expressed as:

$$\ell_{\text{ramp}}(y, f(x)) = \min(1, \max(0, 1 - yf(x))) \quad (3-11)$$

The ramp loss has several important properties:

- 1) Boundedness: Clearly, $0 \leq \ell_{\text{ramp}} \leq 1$ for all inputs, satisfying Eq. (3-9) with $C = 1$.
- 2) Margin interpretation: Like the hinge loss, the ramp loss encourages margin maximization for examples between -1 and 1.
- 3) Saturation effect: Examples with margin less than -1 receive constant penalty, preventing excessive focus on obvious outliers.
- 4) Non-convexity: The ramp loss is non-convex due to the upper truncation, which necessitates careful optimization but contributes to its robustness.

The ramp loss exhibits the redescending influence property characteristic of robust estimators: as $yf(x) \rightarrow -\infty$, the loss saturates at 1, and thus its influence on the risk gradient actually decreases for increasingly misclassified examples. This counterintuitive property is precisely what provides robustness to label noise.

While the ramp loss provides strong robustness, its non-convexity and non-smoothness can complicate optimization. The truncated logistic loss addresses these concerns by applying boundedness to the smooth logistic loss. It is defined as:

$$\ell_{\text{trunc-log}}(y, f(x); \tau) = \min(\log(1 + e^{-yf(x)}), \tau) \quad (3-12)$$

where $\tau > 0$ is a truncation threshold that controls the maximum loss value. In our experiments, we typically use $\tau = 2$, which corresponds to truncating at approximately the 88th percentile of the standard logistic loss distribution.

The truncated logistic loss maintains several desirable properties:

1) Smoothness: Unlike the ramp loss, $\ell_{\text{trunc-log}}$ is smooth (continuously differentiable) everywhere except at the truncation point, facilitating gradient-based optimization.

2) Probabilistic interpretation: For loss values below τ , the truncated logistic loss behaves identically to standard logistic loss, maintaining a natural probabilistic interpretation through the sigmoid function.

3) Adaptive robustness: The parameter τ provides explicit control over the robustness-optimization tradeoff. Smaller τ values increase robustness by more aggressively truncating outliers, while larger τ values preserve more of the standard logistic loss's behavior.

The gradient of the truncated logistic loss is particularly instructive:

$$\nabla_f \ell_{\text{trunc-log}}(y, f(x); \tau) = \begin{cases} -\frac{y}{1 + e^{yf(x)}} & \text{if } \log(1 + e^{-yf(x)}) < \tau \\ 0 & \text{otherwise} \end{cases} \quad (3-13)$$

This gradient reveals the mechanism of robustness: when an example is so severely misclassified that its loss would exceed τ , the gradient becomes exactly zero. The optimizer effectively ignores such examples, preventing them from distorting the learning process.

Both loss functions satisfy the following general properties essential for robust PU learning:

1) Upper boundedness: $\ell(y, f(x)) \leq C$ for all inputs, with $C = 1$ for ramp loss and $C = \tau$ for truncated logistic loss.

2) Classification calibration: Both losses are classification-calibrated, meaning that minimizing the expected risk leads to the Bayes optimal classifier under clean data.

3) Lipschitz continuity: Both losses are Lipschitz continuous with respect to $f(x)$, with constants $L_{\text{ramp}} = 1$ and $L_{\text{trunc-log}} = 1$ (below truncation).

4) Non-negativity: $\ell(y, f(x)) \geq 0$ for all inputs, ensuring compatibility with the non-negative risk formulation.

These properties ensure that our bounded loss formulations maintain the theoretical advantages of standard surrogate losses while gaining robustness to label noise.

3.3.3 Robust PU Risk Estimator

We now integrate bounded loss functions into the PU learning framework to create our robust risk estimator. Starting from the nnPU formulation in Eq. (3-4), we replace the generic loss ℓ with a bounded loss ℓ_b , yielding:

$$\begin{aligned} R_{\text{robust}}(f) &= \pi \cdot \mathbb{E}_{x \sim \tilde{p}(x|Y=+1)}[\ell_b(+1, f(x))] \\ &\quad + \max(0, \mathbb{E}_{x \sim p(x)}[\ell_b(-1, f(x))] \\ &\quad - \pi \cdot \mathbb{E}_{x \sim \tilde{p}(x|Y=+1)}[\ell_b(-1, f(x))]) \end{aligned} \quad (3-14)$$

This formulation assumes access to the noisy positive distribution $\tilde{p}(x | Y = +1)$ as defined in Eq. (3-5). In practice, we only observe samples from this distribution rather than the distribution itself.

The empirical version, used in actual optimization, is:

$$\begin{aligned} \hat{R}_{\text{robust}}(f) &= \frac{\pi}{n_p} \sum_{i=1}^{n_p} \ell_b(+1, f(x_i^p)) \\ &\quad + \max\left(0, \frac{1}{n_u} \sum_{i=1}^{n_u} \ell_b(-1, f(x_i^u))\right. \\ &\quad \left. - \frac{\pi}{n_p} \sum_{i=1}^{n_p} \ell_b(-1, f(x_i^p))\right) \end{aligned} \quad (3-15)$$

where $\{x_i^p\}_{i=1}^{n_p}$ are samples from the (noisy) positive set and $\{x_i^u\}_{i=1}^{n_u}$ are samples from the unlabeled set.

The boundedness of ℓ_b ensures several crucial stability properties for this estimator:

1) Numerical stability: Each term in Eq. (3-15) is bounded, preventing catastrophic cancellation when subtracting large numbers.

2) Gradient boundedness: The gradient of $\hat{R}_{\text{robust}}$ with respect to f is bounded, leading to stable optimization.

3) Controlled variance: The variance of the estimator is naturally controlled by the loss bound C .

4) Robustness to outliers: Mislabeled examples in the positive set cannot dominate the risk estimate due to the loss bound.

The interaction between boundedness and the max operator deserves special attention. The $\max(0, \cdot)$ operation in nnPU already provides some robustness by preventing negative risk estimates. However, this alone is insufficient against label noise because large positive losses

from mislabeled examples can still distort the risk. Boundedness complements this rectification by ensuring that no individual example can produce an excessively large loss.

Our robust estimator can be seen as a generalization that includes standard nnPU as a special case. If we take ℓ_b to be an unbounded loss and let its bound approach infinity, we recover the standard nnPU estimator. Thus, our framework extends rather than replaces existing methodology, providing a smooth pathway to robustness through appropriate loss function choice. The complete workflow of our proposed framework is visually summarized in Figure 3-1, which illustrates the data flow from input partitions through bounded loss application, risk component computation, rectification, and iterative optimization.




Figure 3-1 Architecture of the proposed bounded-loss PU learning framework (Source: Created by author)

3.4 Theoretical Analysis

3.4.1 Robustness Guarantees Under Label Noise

We now establish formal robustness guarantees for our bounded loss PU learning framework. These guarantees characterize how label noise affects learning and demonstrate the advantages of bounded losses.

Theorem 1 (Bounded Risk Distortion). Let ℓ_b be a bounded loss function with $\ell_b(y, f(x)) \leq C$ for all x, y, f . Under the noise model from Eq. (3-5) with noise rate η , the difference between the observed risk and true risk is bounded by:

$$|\tilde{R}_{\text{robust}}(f) - R_{\text{robust}}(f)| \leq 2\pi\eta C \forall f \in \mathcal{F} \quad (3-16)$$

Proof sketch: The difference arises from replacing expectations over $p(x | Y = +1)$ with expectations over $\tilde{p}(x | Y = +1)$. Using the definition $\tilde{p}(x | Y = +1) = (1 - \eta)p(x | Y = +1) + \eta p(x | Y = -1)$, the change in expectation for any function $g(x)$ is:

$$\mathbb{E}_{\tilde{p}}[g(x)] - \mathbb{E}_p[g(x)] = \eta(\mathbb{E}_{p(x|Y=-1)}[g(x)] - \mathbb{E}_{p(x|Y=+1)}[g(x)])$$

Since $|g(x)| \leq C$ when g is a bounded loss, the difference is bounded by $2\eta C$. Accounting for the π multiplier and both occurrences in the risk estimator gives the bound $2\pi\eta C$. $\square$

This theorem has important implications. First, it establishes that the impact of label noise is linear in both the noise rate η and the loss bound C . Second, it demonstrates that bounded losses (with small C) naturally limit noise sensitivity. Third, it provides a practical guideline: for a given expected noise level, we can choose a loss bound C that guarantees risk distortion below a desired threshold.

Corollary 1.1 (Excess Risk Bound). Let $f^* = \arg \min_f R_{\text{robust}}(f)$ be the optimal clean classifier and $\tilde{f}^* = \arg \min_f \tilde{R}_{\text{robust}}(f)$ be the optimal noisy classifier. Under the conditions of Theorem 1:

$$R_{\text{robust}}(\tilde{f}^*) - R_{\text{robust}}(f^*) \leq 4\pi\eta C \quad (3-17)$$

This corollary quantifies the performance degradation due to label noise. Unlike with unbounded losses, where the excess risk can be arbitrarily large, bounded losses guarantee that the suboptimality of the noisy solution is controlled.

Theorem 2 (Uniform Convergence). Let $\mathcal{F}$ be a hypothesis class with Rademacher complexity $\mathfrak{R}_n(\mathcal{F})$. For a bounded loss ℓ_b with bound C and Lipschitz constant L , with probability at least $1 - \delta$:

$$\sup_{f \in \mathcal{F}} |\hat{R}_{\text{robust}}(f) - R_{\text{robust}}(f)| \leq 2\pi L \mathfrak{R}_{n_p}(\mathcal{F}) + 2L \mathfrak{R}_{n_u}(\mathcal{F}) + C \sqrt{\frac{\log(2/\delta)}{2} \left(\frac{\pi^2}{n_p} + \frac{1}{n_u} \right)} \quad (3-18)$$

Proof sketch: This follows from standard Rademacher complexity analysis, noting that the composition of a Lipschitz loss with hypothesis class $\mathcal{F}$ preserves complexity bounds. The weighted combination of samples from positive and unlabeled distributions yields the mixed norm. $\square$

This generalization bound demonstrates that bounded losses lead to tighter concentration inequalities compared to unbounded losses, where the third term would involve unbounded variance.

3.4.2 Bounded Loss Properties and Convergence

The optimization properties of bounded loss functions play a crucial role in the practical success of our framework. We now analyze these properties and establish convergence guarantees.

Definition 1 (Bounded Gradient Property). A loss function ℓ_b has the bounded gradient property if there exists $G < \infty$ such that:

$$\|\nabla_f \ell_b(y, f(x))\| \leq G \forall x, y, f \quad (3-19)$$

Both ramp loss and truncated logistic loss satisfy this property. For ramp loss, $G = 1$; for truncated logistic loss, $G = 1$ for untruncated examples and $G = 0$ for truncated examples.

Theorem 3 (Stochastic Gradient Descent Convergence). Consider minimizing $\hat{R}_{\text{robust}}(f)$ using stochastic gradient descent with step size $\alpha_t = O(1/\sqrt{t})$. Under the bounded gradient property and assuming the risk is μ -strongly convex near the optimum, after T iterations:

$$\mathbb{E}[R_{\text{robust}}(f_T) - R_{\text{robust}}(f^*)] \leq O\left(\frac{G^2}{\mu\sqrt{T}}\right) \quad (3-20)$$

where f_T is the iterate after T steps and f^* is the minimizer.

Proof sketch: Standard SGD analysis for strongly convex functions with bounded gradients gives this rate. The bounded gradient property ensures that the stochastic gradient variance is controlled. $\square$

The bounded gradient property is particularly valuable in the PU learning context due to the subtraction in the risk estimator. In standard nnPU with unbounded losses, the difference $\ell(-1, f(x_u)) - \pi\ell(-1, f(x_p))$ can have unbounded variance even if individual terms have bounded variance. With bounded losses, both terms and their difference have bounded variance, leading to more stable optimization.

Proposition 1 (Noise-Induced Non-Convexity Mitigation). Let $\ell_{\text{ramp}}(y, f(x)) = \min(1, \ell_{\text{hinge}}(y, f(x)))$. For any fixed f , the difference between the gradient on noisy data and clean data is bounded by:

$$\|\nabla_f \tilde{R}_{\text{robust}}(f) - \nabla_f R_{\text{robust}}(f)\| \leq 2\pi\eta G \quad (3-21)$$

This proposition demonstrates that bounded gradients limit the distortion of the optimization landscape due to label noise. While non-convexity may still arise from the loss itself (as with ramp loss), the noise does not introduce additional optimization challenges beyond a bounded perturbation.

Theorem 4 (Population Risk Convergence with Noise). Let f_t be the sequence generated by optimizing $\tilde{R}_{\text{robust}}$ on noisy data. Under appropriate regularity conditions:

$$R_{\text{robust}}(f_t) - R_{\text{robust}}(f^*) \leq O\left(\frac{1}{\sqrt{t}} + \pi\eta C\right) \quad (3-22)$$

This theorem separates the convergence rate into two components: the optimization error decaying as $O(1/\sqrt{t})$ and the statistical error due to noise scaling as $O(\pi\eta C)$. For unbounded losses, the second term would be unbounded, potentially dominating the optimization error regardless of computational budget.

The theoretical analysis establishes that our bounded loss framework provides:

- 1) Controlled risk distortion under label noise (Theorem 1)
- 2) Tight generalization bounds (Theorem 2)
- 3) Stable optimization with convergence guarantees (Theorem 3)
- 4) Explicit error decomposition separating optimization and noise effects (Theorem 4)

These properties collectively justify the bounded loss approach for robust PU learning and provide guidance for practical implementation, such as selecting appropriate loss bounds based on expected noise levels.

3.5 Algorithm Description

3.5.1 Training Procedure (Algorithm 1)

The training procedure for our bounded-loss PU learning framework follows the risk decomposition principle outlined in Eq. (3-15). The algorithm implements stochastic optimization of the robust risk estimator using minibatch gradient descent. The complete training procedure is detailed in Algorithm-1.

 Algorithm-1 Bounded loss positive-unlabeled learning

Step	Description
Input	Labeled positive set $P = \{x_i^p\}_{i=1}^{n_p}$; Unlabeled set $U = \{x_i^u\}_{i=1}^{n_u}$; Class prior $\pi \in (0,1)$; Learning rate $\alpha > 0$; Number of epochs T ; Model parameters θ (initialized); Bounded loss $\ell_b \in \{\ell_{\text{ramp}}, \ell_{\text{trunc-log}}\}$; Batch sizes B_p, B_u
Output	Trained classifier $f_\theta: \mathcal{X} \rightarrow \mathbb{R}$
1	for epoch $t = 1$ to T do
2	Shuffle P and U randomly
3	for each minibatch ($P_{\text{batch}} \subseteq P, U_{\text{batch}} \subseteq U$) do
4	Compute positive loss (positive label):
5	$L_{\text{pos1}} = \frac{\pi}{ P_{\text{batch}} } \sum_{x \in P_{\text{batch}}} \ell_b(+1, f_\theta(x))$
6	Compute positive loss (negative label):
7	$L_{\text{pos2}} = \frac{\pi}{ P_{\text{batch}} } \sum_{x \in P_{\text{batch}}} \ell_b(-1, f_\theta(x))$
8	Compute unlabeled loss:
9	$L_{\text{unlab}} = \frac{1}{ U_{\text{batch}} } \sum_{x \in U_{\text{batch}}} \ell_b(-1, f_\theta(x))$
10	Apply non-negative risk rectification:
11	$L_{\text{neg}} = \max(0, L_{\text{unlab}} - L_{\text{pos2}})$
12	Compute total empirical risk:
13	$R_{\text{total}} = L_{\text{pos1}} + L_{\text{neg}}$
14	Update model parameters:
15	$\theta \leftarrow \theta - \alpha \nabla_\theta R_{\text{total}}$
16	end for
17	end for
18	return f_θ

(Source: Created by author)

Algorithm-1 outlines the complete training procedure for our bounded-loss PU learning framework. The algorithm follows a standard stochastic optimization scheme with several key components specific to PU learning: (1) separate processing of positive and unlabeled batches, (2) computation of both positive and negative risk components using the bounded loss function ℓ_b , (3) application of the non-negative rectification operator $\max(0, \cdot)$ to ensure stable risk estimates, and (4) aggregation of risk components for gradient computation. The algorithm maintains computational efficiency comparable to standard supervised learning while incorporating the necessary adjustments for PU learning and bounded loss optimization.

3.5.2 Implementation Details

The practical implementation of Algorithm 1 involves several important considerations to ensure efficiency, stability, and reproducibility. Our implementation uses Python with PyTorch for automatic differentiation and GPU acceleration. Key implementation details include:

1) **Model Architecture:** We employ a fully-connected neural network with ReLU activations and batch normalization for tabular data, and a convolutional neural network for image data. The final layer uses a linear activation without sigmoid, as the bounded losses operate directly on logits.

2) **Optimization:** We use the Adam optimizer with default parameters ($\beta_1 = 0.9$, $\beta_2 = 0.999$, $\varepsilon = 10^{-8}$). The learning rate follows a cosine annealing schedule, starting at $\alpha = 0.001$ and decaying to $\alpha = 0.00001$ over T epochs.

3) **Batch Composition:** We maintain separate batch loaders for positive and unlabeled data with batch sizes $B_p = \min(64, n_p)$ and $B_u = \min(256, n_u)$, ensuring computational efficiency while preserving statistical properties.

4) **Loss Function Implementation:** The ramp loss is implemented using `torch.clamp()` for efficient truncation. The truncated logistic loss uses a conditional assignment to apply truncation when $\log(1 + \exp(-y f(x))) > \tau$.

5) **Numerical Stability:** We add $\varepsilon = 10^{-8}$ to denominators and logarithmic arguments to prevent numerical underflow. Gradient clipping with threshold 1.0 is applied to maintain stability during training.

6) **Early Stopping:** Training terminates if the validation loss (computed on a held-out subset) does not improve for 20 consecutive epochs, with patience restored upon improvement.

7) **Reproducibility:** We fix random seeds for PyTorch, NumPy, and Python's random module to ensure reproducible results. All hyperparameters are logged using Weights & Biases for experiment tracking.

3.6 Evaluation Metrics

3.6.1 Accuracy, Precision, Recall, F1-Score

We employ standard classification metrics to evaluate model performance. For binary classification with predicted labels $\hat{y} \in \{0,1\}$ and true labels $y \in \{0,1\}$, we define:

1) **Accuracy:** Proportion of correctly classified samples:

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN} \quad (3-23)$$

2) Precision: Proportion of true positives among predicted positives:

$$\text{Precision} = \frac{TP}{TP + FP} \quad (3-24)$$

3) Recall (Sensitivity): Proportion of true positives among actual positives:

$$\text{Recall} = \frac{TP}{TP + FN} \quad (3-25)$$

4) F1-Score: Harmonic mean of precision and recall:

$$F1 = 2 \times \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}} \quad (3-26)$$

where TP = True Positives, TN = True Negatives, FP = False Positives, FN = False Negatives.

3.6.2 ROC-AUC and PR-AUC

For probabilistic evaluation and class-imbalance scenarios, we use area-under-curve metrics:

1) ROC-AUC: Area under the Receiver Operating Characteristic curve, plotting True Positive Rate (Recall) against False Positive Rate ($FPR = FP/(FP+TN)$) across classification thresholds. ROC-AUC measures overall ranking ability independent of threshold choice.

2) PR-AUC: Area under the Precision-Recall curve, particularly informative for imbalanced datasets where negative class dominates. PR-AUC focuses on positive class performance and is more sensitive to class distribution changes.

Both metrics are computed using trapezoidal integration with threshold values sampled at 100 equally spaced points between 0 and 1.

3.6.3 Training Stability Metrics

To evaluate optimization stability and robustness properties, we define additional metrics:

1) Loss Variance: Variance of training loss across the final 10% of epochs, measuring convergence stability:

$$\sigma_L^2 = \frac{1}{N} \sum_{i=1}^N (L_i - \bar{L})^2 \quad (3-27)$$

2) Gradient Norm Statistics: Mean and maximum ℓ_2 -norm of parameter gradients during training, indicating optimization smoothness:

$$\bar{g} = \frac{1}{T} \sum_{t=1}^T \|\nabla_{\theta} R_t\|_2 \quad (3-28)$$

3) Noise Sensitivity Ratio: Performance degradation ratio between clean and noisy conditions:

$$\text{NSR} = \frac{\text{Perf}_{\text{noisy}}}{\text{Perf}_{\text{clean}}} \quad (3-29)$$

where Perf can be any primary metric (Accuracy, ROC-AUC, etc.).

4) Convergence Speed: Number of epochs required to reach 95% of final performance, measuring optimization efficiency.

These comprehensive metrics provide multi-dimensional evaluation of both classification performance and training characteristics, enabling thorough comparison of robustness properties across methods and noise conditions.

3.7 Chapter Summary

This chapter presents the methodological foundation for a robust Positive-Unlabeled learning framework that integrates bounded loss functions to address label noise in the positive set. The methodology is structured to achieve theoretical robustness, algorithmic efficiency, and practical applicability.

Section 3.1 introduces the core premise: label noise in PU learning requires integrated solutions at the algorithmic level rather than external preprocessing. The approach reconceptualizes noise as an intrinsic property addressed through loss function design and risk estimation, guided by principles of theoretical soundness, practical feasibility, and robustness by design.

Section 3.2 establishes the formal problem setup. PU learning involves labeled positive examples drawn from the true positive distribution and unlabeled examples from a marginal distribution mixing both classes. The unbiased PU (uPU) risk estimator provides theoretical elegance but suffers from high variance, while the non-negative PU (nnPU) estimator introduces a rectification operator for stability. A label noise model is formalized where the observed positive distribution is a mixture of true positives and mislabeled negatives with noise rate η . This unidirectional contamination creates risk distortion that grows with both noise rate and loss function unboundedness.

Section 3.3 develops the proposed bounded loss framework. The motivation stems from the pathological dynamics of unbounded losses, where mislabeled positives dominate risk estimation through arbitrarily large penalties. Bounded losses enforce a maximum contribution per example, providing natural noise robustness, stabilized risk decomposition, and improved gradient dynamics. Two specific formulations are detailed: the ramp loss, which truncates hinge loss at 1, offering strong robustness through its redescending influence property; and the truncated logistic loss, which caps standard logistic loss at threshold τ , maintaining smoothness and tunable robustness. Both are integrated into the nnPU framework to create a robust risk estimator that controls variance, stabilizes optimization, and limits outlier influence.

Section 3.4 provides theoretical analysis establishing formal guarantees. Theorem 1 bounds risk distortion under label noise by $2\pi\eta C$, where C is the loss bound, demonstrating linear dependence on noise rate and loss bound. Corollary 1.1 extends this to excess risk bounds. Theorem 2 establishes generalization bounds through Rademacher complexity. The bounded gradient property ensures stable SGD convergence with rate $O(1/\sqrt{T})$, while Theorem 4 decomposes error into optimization and noise components, showing that bounded losses prevent unbounded noise-induced degradation.

Section 3.5 presents the algorithmic implementation through Algorithm-1, which performs minibatch stochastic optimization of the robust risk estimator. The procedure separately processes positive and unlabeled batches, computes both risk components using bounded losses, applies non-negative rectification, and updates parameters via gradient descent. Implementation details cover model architecture, optimization scheduling, loss function implementation, and reproducibility measures.

Section 3.6 defines comprehensive evaluation metrics spanning classification performance (accuracy, precision, recall, F1-score), probabilistic assessment (ROC-AUC, PR-AUC), and training stability (loss variance, gradient norms, noise sensitivity ratio, convergence speed). These metrics enable multi-dimensional comparison of robustness properties across methods and noise conditions.

Chapter IV EXPERIMENT RESULTS AND DISCUSSIONS

4.1 Experimental Setup

This chapter details the comprehensive experimental validation of the proposed robust Positive-Unlabeled learning framework. Through systematic evaluation across diverse datasets and noise conditions, we aim to empirically validate the theoretical claims established in Chapter 3 and demonstrate the practical advantages of bounded loss functions in noisy PU learning scenarios. The experimental design follows principles of reproducibility, fairness in comparison, and comprehensive evaluation across multiple performance dimensions.

4.1.1 Dataset Description (Synthetic, UCI Heart, MNIST-PU)

To ensure robust evaluation across different data characteristics and application domains, we employ three distinct datasets: one synthetic dataset for controlled analysis, one real-world medical dataset for practical validation, and one image dataset for high-dimensional scenario testing.

1) Synthetic Dataset: We generate a two-class Gaussian mixture dataset with controllable overlap and noise injection capability. The positive class follows $\mathcal{N}(\mu_+, \Sigma_+)$ with $\mu_+ = [1.5, 1.5]^T$ and $\Sigma_+ = \text{diag}(0.8, 0.8)$, while the negative class follows $\mathcal{N}(\mu_-, \Sigma_-)$ with $\mu_- = [-1.5, -1.5]^T$ and $\Sigma_- = \text{diag}(0.8, 0.8)$. The dataset comprises 2,000 samples (1,000 per class) with 25 features, including 22 noise features to increase difficulty. The class prior is set to $\pi = 0.5$ for balanced conditions. This synthetic setup allows precise control over noise injection and distribution overlap, enabling isolation of algorithm behaviors under varying contamination levels.

2) UCI Heart Disease Dataset: This real-world medical dataset contains 303 clinical instances with 13 diagnostic features for predicting presence of heart disease. We treat confirmed heart disease cases as positives ($\approx 45\%$ of data) and the remaining as unlabeled data following standard PU protocol. The dataset presents realistic challenges including feature heterogeneity, moderate class imbalance, and inherent diagnostic uncertainty. Feature preprocessing includes standardization and handling of missing values via median imputation. The dataset's medical relevance makes it particularly suitable for evaluating robustness in high-stakes applications where label noise is prevalent.

3) MNIST-PU Dataset: We adapt the standard MNIST handwritten digit dataset to a PU learning scenario by treating digit '0' as the positive class and digits '1' through '9' as unlabeled data. From the 70,000 original images, we create a subset of 10,000 samples with balanced sampling across digits. Each 28×28 grayscale image is flattened to 784-dimensional feature vectors. This dataset introduces high-dimensionality challenges and tests the framework's scalability to image data while simulating practical applications like rare object detection or novelty identification. Table 4-1 summarizes the dataset properties and experimental setup.

Table 4-1 Dataset characteristics and experimental configuration

Dataset	Type	Samples	Features	Positive Ratio (π)	Train/Test Split	Primary Evaluation Focus
Synthetic	Artificial Gaussian	2,000	25 (2 signal + 22 noise)	0.5	70%/30%	Controlled noise analysis; Fundamental behavior
UCI Heart	Medical tabular	303	13 clinical features	0.45	70%/30% stratified	Real-world robustness; Medical application relevance
UCI Heart	Medical tabular	303	13 clinical features	0.45	70%/30% stratified	Real-world robustness; Medical application relevance
MNIST-PU	Image pixels	10,000	784 (28×28 flattened)	0.1	80%/20%	High-dimensional scalability; Image classification

(Source: Created by author)

4.1.2 Hardware and Software Configurations

All experiments were conducted on a standardized computing environment to ensure reproducibility and fair comparison between methods. The hardware configuration consisted of an Intel Core i9-12900K CPU (16 cores, 24 threads, 3.2-5.2 GHz), an NVIDIA RTX 4090 GPU with 24GB GDDR6X memory, 64GB DDR5 system RAM at 4800 MHz, and 2TB of NVMe PCIe 4.0 SSD storage.

The software environment used Ubuntu 22.04 LTS as the operating system with Python version 3.9.18. The deep learning framework was PyTorch 2.0.1 with CUDA 11.8 support. Key libraries included NumPy 1.24.3, SciPy 1.10.1, and scikit-learn 1.3.0, while visualization utilized Matplotlib 3.7.1 and Seaborn 0.12.2.

Regarding implementation specifics, all neural network models were implemented using PyTorch with automatic differentiation. Optimization employed the Adam optimizer with parameters $\beta_1=0.9$, $\beta_2=0.999$, and $\epsilon=10^{-8}$. Batch sizes were set to 64 for positive samples and 256 for unlabeled samples. The learning rate began at 0.001 with cosine annealing to 0.00001 over the training duration. Early stopping was implemented with a patience of 20 epochs based

on validation loss. Random seeds were fixed across all experiments: 42 for data splitting, 123 for model initialization, and 456 for noise injection.

From a computational perspective, each experimental run (method $\times$ dataset $\times$ noise level) required approximately 15 to 45 minutes depending on dataset size and model complexity, with total experimental time exceeding 200 hours for comprehensive evaluation across all configurations. Memory usage remained below 12GB for all experiments, while GPU utilization averaged between 85% and 95% during training.

4.1.3 Noise Injection Protocol

To systematically evaluate robustness under label contamination, we implement a controlled noise injection protocol that simulates realistic labeling errors in the positive set. The protocol follows the noise model defined in Eq. (3-5) with specified noise rates.

Noise Injection Procedure

For each dataset and target noise rate $\eta \in \{0.0, 0.2, 0.4, 0.6\}$, we execute the following contamination process:

- 1) True Positive Identification: From the complete dataset with ground truth labels, identify all true positive instances $\mathcal{P}_{\text{true}} = \{x: y = +1\}$.
- 2) True Negative Pool Construction: Identify all true negative instances $\mathcal{N}_{\text{true}} = \{x: y = -1\}$.
- 3) Contamination Sample Selection: Randomly select $n_{\text{noise}} = \lceil \eta \times |\mathcal{P}_{\text{true}}| \rceil$ instances from $\mathcal{N}_{\text{true}}$ without replacement.
- 4) Label Corruption: Assign positive labels to the selected negative instances, creating corrupted positives $\mathcal{P}_{\text{corrupt}}$.
- 5) Final Positive Set Construction: Combine $\mathcal{P}_{\text{clean}} = \mathcal{P}_{\text{true}} \setminus \mathcal{P}_{\text{corrupt}}$ (with $|\mathcal{P}_{\text{clean}}| = (1 - \eta) \times |\mathcal{P}_{\text{true}}|$) and $\mathcal{P}_{\text{corrupt}}$ to form the noisy positive set $\mathcal{P}_{\text{noisy}}$.
- 6) Unlabeled Set Construction: The unlabeled set comprises all remaining instances: $\mathcal{U} = (\mathcal{N}_{\text{true}} \setminus \mathcal{P}_{\text{corrupt}}) \cup \mathcal{P}_{\text{corrupt}}$ (though the latter are now labeled as positives in $\mathcal{P}_{\text{noisy}}$).

Noise Characteristics

- 1) Uniform Random Noise: The selection in step 3 follows uniform random sampling, implementing the symmetric noise assumption.
- 2) Feature-Independent: Noise injection is independent of instance features, ensuring controlled experimental conditions.

3) Reproducible: Random seeds ensure identical noise patterns across different method evaluations for fair comparison.

4) Validation of Noise Injection: We verify the noise injection process by measuring the effective noise rate in the positive set:

$$\eta_{\text{effective}} = \frac{|\mathcal{P}_{\text{corrupt}}|}{|\mathcal{P}_{\text{noisy}}|} \approx \eta \quad (4 - 1)$$

with observed deviations below 1% across all experiments due to rounding to discrete sample counts. Table 4-2 shows the noise level specifications and corresponding sample counts.

Table 4-2 Noise level specifications and corresponding sample counts

Noise Rate (η)	Synthetic Dataset	UCI Heart Dataset	MNIST-PU Dataset
0.0 (Clean)	500 clean positives	136 clean positives	800 clean positives
0.2 (Low Noise)	400 clean + 100 noisy	109 clean + 27 noisy	640 clean + 160 noisy
0.4 (Medium Noise)	300 clean + 200 noisy	82 clean + 54 noisy	480 clean + 320 noisy
0.6 (High Noise)	200 clean + 300 noisy	54 clean + 82 noisy	320 clean + 480 noisy

(Source: Created by author)

4.2 Performance Under Varying Noise Levels

4.2.1 Quantitative Results on Synthetic Data

The synthetic dataset provides a controlled environment to isolate and analyze the impact of label noise on learning performance. Table 4-3 presents comprehensive results comparing our proposed bounded-loss methods against baseline approaches.

Table 4.3 Performance comparison on synthetic dataset across noise levels

Noise Rate (η)	Method	Accuracy (%)	Precision (%)	Recall (%)	F1-Score (%)	ROC-AUC (%)	PR-AUC (%)
0.0	uPU	85.2 ± 1.3	84.8 ± 1.5	85.1 ± 1.4	84.9 ± 1.4	87.5 ± 1.2	86.2 ± 1.3
	nnPU	87.4 ± 1.1	86.5 ± 1.2	87.8 ± 1.1	87.1 ± 1.1	89.6 ± 0.9	88.4 ± 1.0
	Ramp-Loss	89.3 ± 0.9	88.9 ± 1.0	89.7 ± 0.8	89.3 ± 0.9	92.1 ± 0.7	91.0 ± 0.8
	Trunc-Log	88.2 ± 1.0	87.8 ± 1.1	88.5 ± 1.0	88.1 ± 1.0	90.8 ± 0.8	89.7 ± 0.9
0.2	uPU	77.1 ± 1.8	75.3 ± 2.0	77.8 ± 1.9	76.5 ± 1.9	79.6 ± 1.6	77.8 ± 1.7
	nnPU	81.2 ± 1.5	79.8 ± 1.7	82.1 ± 1.5	80.9 ± 1.6	83.4 ± 1.3	82.1 ± 1.4
	Ramp-Loss	84.4 ± 1.2	83.2 ± 1.3	85.0 ± 1.1	84.1 ± 1.2	87.2 ± 1.0	86.0 ± 1.1
	Trunc-Log	83.1 ± 1.3	82.0 ± 1.4	83.7 ± 1.2	82.8 ± 1.3	85.9 ± 1.1	84.8 ± 1.2
0.4	uPU	66.3 ± 2.4	64.1 ± 2.6	67.0 ± 2.5	65.5 ± 2.5	69.5 ± 2.2	66.8 ± 2.3
	nnPU	69.8 ± 2.1	67.9 ± 2.3	70.5 ± 2.2	69.2 ± 2.2	72.9 ± 1.9	70.6 ± 2.0
	Ramp-Loss	78.2 ± 1.6	76.8 ± 1.8	79.1 ± 1.6	77.9 ± 1.7	81.7 ± 1.4	80.3 ± 1.5
	Trunc-Log	76.4 ± 1.7	75.0 ± 1.9	77.3 ± 1.7	76.1 ± 1.8	79.8 ± 1.5	78.5 ± 1.6
0.6	uPU	52.1 ± 3.1	50.4 ± 3.3	53.8 ± 3.2	52.0 ± 3.2	57.5 ± 2.8	53.2 ± 2.9
	nnPU	56.3 ± 2.8	54.7 ± 3.0	57.5 ± 2.9	56.0 ± 2.9	60.8 ± 2.5	57.6 ± 2.6
	Ramp-Loss	69.1 ± 2.0	68.0 ± 2.2	70.2 ± 2.0	69.1 ± 2.1	73.8 ± 1.8	72.3 ± 1.9
	Trunc-Log	67.2 ± 2.1	66.1 ± 2.3	68.3 ± 2.1	67.2 ± 2.2	71.8 ± 1.9	70.2 ± 2.0

(Source: Created by author)

As shown in Table 4-3, our proposed bounded-loss methods consistently Outperform baseline approaches across all noise levels and metrics. The performance advantage becomes more pronounced as noise increases, with the ramp-loss variant showing particularly strong robustness. At 60% noise, our ramp-loss method maintains 69.1% accuracy compared to 56.3% for nnPU and 52.1% for uPU, representing a 22.7% and 32.6% relative improvement respectively.




Figure 4-1 Accuracy degradation with increasing noise (Source: Created by author)

The accuracy degradation pattern (Figure 4-1) reveals several important insights. First, all methods show performance decline as noise increases, as expected. Second, the bounded-loss methods (green and magenta curves) exhibit shallower degradation slopes, indicating better noise resistance. Third, the divergence between methods becomes more pronounced at higher noise levels, suggesting that bounded losses provide increasing relative advantages in highly contaminated environments. The 22.7% improvement over nnPU and 32.6% improvement over uPU at $\eta=0.6$ (highlighted in the figure annotation) demonstrate the substantial practical benefits of our approach in challenging conditions.

4.2.2 Results on Real-World Datasets (UCI, MNIST)

Table 4-4 presents performance results on real-world datasets, confirming that the advantages observed on synthetic data generalize to practical applications with realistic data characteristics.

Table 4-4 Performance on real-world datasets under varying noise levels

Dataset	Noise Rate	Method	Accuracy (%)	ROC-AUC (%)	PR-AUC (%)	Training Stability ($\sigma^2 \times 10^3$)
UCI Heart	0.0	uPU	87.2 ± 1.8	88.7 ± 1.6	87.5 ± 1.7	5.2 ± 0.8
		nnPU	89.1 ± 1.5	90.3 ± 1.3	89.4 ± 1.4	4.1 ± 0.6
		Ramp-Loss	92.3 ± 1.2	93.1 ± 1.0	92.8 ± 1.1	2.3 ± 0.4
		Trunc-Log	91.4 ± 1.3	92.2 ± 1.1	91.6 ± 1.2	2.8 ± 0.5
	0.2	uPU	78.5 ± 2.1	80.1 ± 1.9	78.8 ± 2.0	8.7 ± 1.2
		nnPU	83.2 ± 1.8	84.9 ± 1.6	83.6 ± 1.7	6.3 ± 0.9
		Ramp-Loss	87.4 ± 1.4	88.8 ± 1.2	87.9 ± 1.3	3.1 ± 0.5
		Trunc-Log	86.2 ± 1.5	87.5 ± 1.3	86.7 ± 1.4	3.5 ± 0.6
	0.4	uPU	70.3 ± 2.6	72.1 ± 2.3	70.8 ± 2.4	12.5 ± 1.7
		nnPU	74.6 ± 2.3	76.4 ± 2.1	75.1 ± 2.2	9.2 ± 1.3
		Ramp-Loss	81.8 ± 1.7	83.2 ± 1.5	82.3 ± 1.6	3.9 ± 0.6
		Trunc-Log	80.3 ± 1.8	81.7 ± 1.6	80.9 ± 1.7	4.3 ± 0.7
	0.6	uPU	58.1 ± 3.2	60.2 ± 2.9	58.7 ± 3.0	18.3 ± 2.4
		nnPU	62.8 ± 2.9	64.6 ± 2.6	63.4 ± 2.7	13.7 ± 1.9
		Ramp-Loss	73.5 ± 2.0	75.1 ± 1.8	74.2 ± 1.9	4.8 ± 0.7
		Trunc-Log	71.6 ± 2.1	73.2 ± 1.9	72.3 ± 2.0	5.2 ± 0.8
MNIST-PU	0.0	uPU	74.3 ± 1.5	77.8 ± 1.3	75.6 ± 1.4	6.8 ± 0.9
		nnPU	77.6 ± 1.3	81.2 ± 1.1	78.9 ± 1.2	5.3 ± 0.7
		Ramp-Loss	82.4 ± 1.0	85.9 ± 0.8	83.7 ± 0.9	3.1 ± 0.4
		Trunc-Log	81.3 ± 1.1	84.7 ± 0.9	82.5 ± 1.0	3.6 ± 0.5
	0.2	uPU	66.7 ± 1.8	71.3 ± 1.6	67.9 ± 1.7	9.5 ± 1.2
		nnPU	71.8 ± 1.6	76.5 ± 1.4	73.2 ± 1.5	7.4 ± 1.0
		Ramp-Loss	77.9 ± 1.2	82.1 ± 1.0	79.3 ± 1.1	3.7 ± 0.5
		Trunc-Log	76.5 ± 1.3	80.8 ± 1.1	77.9 ± 1.2	4.2 ± 0.6
	0.4	uPU	54.8 ± 2.3	61.2 ± 2.1	56.7 ± 2.2	14.3 ± 1.8
		nnPU	60.3 ± 2.1	66.9 ± 1.9	62.4 ± 2.0	11.2 ± 1.5
		Ramp-Loss	71.5 ± 1.5	77.4 ± 1.3	73.2 ± 1.4	4.5 ± 0.6
		Trunc-Log	69.8 ± 1.6	75.6 ± 1.4	71.4 ± 1.5	5.0 ± 0.7
	0.6	uPU	41.2 ± 2.9	49.7 ± 2.6	43.8 ± 2.7	21.8 ± 2.7
		nnPU	47.5 ± 2.6	56.4 ± 2.3	50.3 ± 2.4	17.3 ± 2.2
		Ramp-Loss	62.8 ± 1.8	70.5 ± 1.6	65.4 ± 1.7	5.4 ± 0.7
		Trunc-Log	60.4 ± 1.9	68.1 ± 1.7	63.1 ± 1.8	5.9 ± 0.8

(Source: Created by author)

The results in Table 4-4 and Figure 4-2 demonstrate several consistent patterns across real-world datasets. First, the training stability metric (σ^2) shows dramatically lower values for bounded-loss methods, indicating smoother optimization and reduced sensitivity to noise. At $\eta=0.6$ on UCI Heart, our ramp-loss method achieves $\sigma^2=4.8 \times 10^{-3}$ compared to 18.3×10^{-3} for uPU, a 73.8% reduction in training variance. Second, the ROC-AUC advantage maintains across domains, with our method showing particular strength in high-noise regimes. Third, the MNIST-PU results confirm that bounded losses scale effectively to high-dimensional image data, overcoming challenges that often exacerbate noise sensitivity in complex models.




Figure 4-2 Cross-dataset ROC-AUC comparison (Source: Created by author)

4.3 Comparison with Baseline Models

4.3.1 uPU vs. nnPU vs. Proposed Method

Table 4-5 provides a direct comparative analysis at the critical medium-noise level ($\eta=0.4$), representing realistic contamination levels encountered in practical applications.

Table 4-5 Direct comparison at medium noise level ($\eta=0.4$) across all datasets

Dataset	Metric	uPU	nnPU	Ramp-Loss (Ours)	Improvement over uPU	Improvement over nnPU
Synthetic	Accuracy	66.3	69.8	78.2	+17.9%	+12.0%
	ROC-AUC	69.5	72.9	81.7	+17.6%	+12.1%
	PR-AUC	66.8	70.6	80.3	+20.2%	+13.7%
UCI Heart	Accuracy	70.3	74.6	81.8	+16.4%	+9.7%
	ROC-AUC	72.1	76.4	83.2	+15.4%	+8.9%
	PR-AUC	70.8	75.1	82.3	+16.2%	+9.6%
MNIST-PU	Accuracy	54.8	60.3	71.5	+30.5%	+18.6%
	ROC-AUC	61.2	66.9	77.4	+26.5%	+15.7%
	PR-AUC	56.7	62.4	73.2	+29.1%	+17.3%
Average	All Metrics	65.4	69.6	78.5	+20.1%	+12.8%

(Source: Created by author)

The comparative analysis reveals consistent advantages across all datasets and metrics. Our ramp-loss method achieves an average improvement of 20.1% over uPU and 12.8% over nnPU. Notably, the improvement is most substantial on the MNIST-PU dataset (30.5% over uPU, 18.6% over nnPU in accuracy), suggesting that bounded losses provide particular benefits in high-dimensional settings where noise effects are amplified.

4.3.2 Performance Trends with Increasing Noise

Table 4-6 Performance degradation analysis from clean to 60% noise

Method	Synthetic Data	UCI Heart	MNIST- PU	Average Degradation	-	-	-	-
	Acc Drop	AUC Drop	Acc Drop	AUC Drop	Acc Drop	AUC Drop	Acc	AUC
uPU	-33.1%	-30.0%	-29.1%	-28.5%	-33.1%	-28.1%	-31.8%	-28.9%
nnPU	-31.1%	-28.8%	-26.3%	-25.7%	-30.1%	-24.8%	-29.2%	-26.4%
Ramp-Loss	-20.2%	-18.3%	-18.8%	-18.0%	-19.6%	-15.4%	-19.5%	-17.2%
Trunc-Log	-21.0%	-19.0%	-19.8%	-19.0%	-20.9%	-16.6%	-20.6%	-18.2%

(Source: Created by author)

The degradation analysis in Table 4-6 reveals important insights about noise resilience. Our proposed methods exhibit significantly shallower performance drops as noise increases: an average accuracy degradation of 19.5% for ramp-loss compared to 31.8% for uPU and 29.2% for nnPU. This represents a 38.7% reduction in degradation rate compared to uPU and 33.2% reduction compared to nnPU.




Figure 4-3 Noise resilience comparison across methods (Source: Created by author)

Figure 4-3 provides a visual comparison of noise resilience using a radar plot, where larger areas indicate better performance retention across datasets. The bounded-loss methods (green and magenta) clearly dominate, covering substantially more area than baseline methods (red and blue). This visualization confirms that bounded losses provide consistent advantages across diverse data types and characteristics.

The analysis demonstrates several key patterns:

1) Consistent Advantage: Bounded losses outperform baselines across all noise levels, datasets, and metrics.

2) Increasing Advantage: The relative improvement grows with noise level, confirming the theoretical prediction that boundedness provides greater benefits in highly contaminated environments.

3) Stability Benefits: Beyond accuracy improvements, bounded losses provide dramatically better training stability (Table 4-4), reducing variance by 70-80% in high-noise conditions.

4) Domain Generalization: Advantages hold across synthetic, medical, and image data, suggesting broad applicability.

These results strongly support our theoretical claims that bounded loss functions provide inherent robustness to label noise in PU learning, offering practical benefits that increase with contamination severity.

4.4 Training Behavior and Convergence Analysis

4.4.1 Loss Curves and Optimization Stability

The optimization dynamics of PU learning methods under label noise reveal critical insights into their stability and convergence properties. Figure 4-4 presents the training loss curves for all methods across 100 training epochs on the synthetic dataset with $\eta=0.4$ noise contamination.




Figure 4-4 Training loss curves (Source: Created by author)

The loss curves in Figure 4-4 reveal striking differences in optimization stability. Traditional methods (uPU and nnPU) exhibit high-frequency oscillations and slow convergence due to their sensitivity to mislabeled examples. In contrast, bounded-loss methods show smooth, monotonic convergence with minimal oscillation. The ramp loss demonstrates particularly rapid convergence, reaching near-optimal loss values by epoch 40, while uPU continues to oscillate even after 100 epochs.

Table 4-7 Quantitative optimization stability metrics

Method	Final Loss	Convergence Epoch*	Loss Variance ($\sigma^2 \times 10^3$)	Gradient Norm (Mean)	Training Time (seconds)
uPU	0.087 ± 0.015	82 ± 12	12.8 ± 2.1	0.45 ± 0.08	142 ± 8
nnPU	0.063 ± 0.009	65 ± 9	8.3 ± 1.5	0.32 ± 0.06	138 ± 7
Ramp-Loss	0.038 ± 0.004	38 ± 5	1.7 ± 0.3	0.18 ± 0.03	121 ± 6
Trunc-Log	0.041 ± 0.005	42 ± 6	2.1 ± 0.4	0.21 ± 0.04	125 ± 6

(Source: Created by author)

The optimization stability metrics in Table 4-7 quantify these observations. Our ramp-loss method achieves 74% lower loss variance and 60% lower average gradient norm compared to nnPU, explaining its smoother convergence. Remarkably, it converges in approximately half the epochs required by baseline methods (38 vs 82 for uPU), translating to 15% reduction in training time despite the non-convex nature of ramp loss.

4.4.2 Sensitivity to Hyperparameters

The truncated logistic loss introduces an important hyperparameter τ controlling the truncation threshold. Figure 4-5 systematically analyzes the sensitivity of model performance to τ across different noise levels. Table 4-8 presents the Optimal τ ranges for different noise conditions.




Figure 4-5 Truncation threshold sensitivity analysis (Source: Created by author)

Table 4-8 Optimal τ ranges for different noise conditions

Noise Level (η)	Optimal τ Range	Peak Accuracy	Robustness-Retention Tradeoff
0.0 (Clean)	2.0 - 3.0	91.4% at $\tau=2.0$	Standard logistic behavior
0.2 (Low)	1.5 - 2.5	86.2% at $\tau=2.0$	Balanced robustness
0.4 (Medium)	1.0 - 2.0	80.3% at $\tau=1.5$	Enhanced robustness
0.6 (High)	0.5 - 1.5	71.6% at $\tau=1.0$	Maximum robustness

(Source: Created by author)

The sensitivity analysis reveals several important patterns:

- 1) Clean Data Performance: Larger τ values (2.0-3.0) perform best under clean conditions, approaching standard logistic loss behavior.
- 2) Noisy Data Performance: Smaller τ values provide better robustness, with optimal τ decreasing as noise increases.
- 3) Robustness-Accuracy Tradeoff: The performance retention metric shows a clear peak around $\tau=1.5$ -2.0, indicating an optimal balance.
- 4) Stability: Lower τ values provide better parameter stability (less sensitivity to τ variations), though extremely small τ (<1.0) can underfit.

Based on this analysis, we recommend $\tau=2.0$ as a general default value, with $\tau=1.5$ for applications expecting moderate noise ($\eta \approx 0.4$) and $\tau=1.0$ for severely noisy environments.

4.5 Ablation Studies

4.5.1 Effect of Bounded Loss vs. Rectification

To disentangle the contributions of bounded loss functions from the non-negative rectification inherent in nnPU, we conduct a systematic ablation study comparing four configurations at $\eta=0.4$ noise level. Table 4-9 and Figure 4-6 demonstrate the ablation study results.

Table 4-9 Ablation study results (synthetic data, $\eta=0.4$)

Configuration	Bounded Loss	Rectification	Accuracy (%)	ROC-AUC (%)	Training Stability ($\sigma^2 \times 10^3$)	Relative to nnPU
A: uPU	No	No	66.3 ± 2.4	69.5 ± 2.2	12.8 ± 2.1	Baseline
B: nnPU	No	Yes	69.8 ± 2.1	72.9 ± 1.9	8.3 ± 1.5	+3.5% accuracy
C: Bounded-only	Yes	No	75.1 ± 1.8	78.4 ± 1.6	3.2 ± 0.6	+8.8% accuracy
D: Full (Ours)	Yes	Yes	78.2 ± 1.6	81.7 ± 1.4	1.7 ± 0.3	+11.9% accuracy

(Source: Created by author)




Figure 4-6 Ablation study visualization (Source: Created by author)

The ablation study reveals crucial insights into component contributions:

1) Bounded Loss Alone: Configuration C (bounded loss without rectification) already provides substantial improvement (+8.8% accuracy over uPU), demonstrating that boundedness itself confers significant robustness.

2) Rectification Alone: Configuration B (rectification without bounded loss) provides moderate improvement (+3.5%), primarily through stabilization of risk estimates.

3) Synergistic Combination: Configuration D (full method) achieves the best performance (+11.9% over uPU), showing that bounded loss and rectification operate through complementary mechanisms to provide robustness.

4) Stability Amplification: The training stability improvement is particularly dramatic, with the full method achieving 588.2 stability index vs 78.1 for uPU, a $7.5\times$ improvement.

4.5.2 Component-Wise Contribution Analysis

We further decompose the risk estimator to analyze how different components contribute to overall performance. Figure 4-7 examines the behavior of the positive risk component (R_{pos}) and negative risk component (R_{neg}) during training. Table 4-10 shows the statistical analysis.




Figure 4-7 Risk component analysis (Source: Created by author)

Table 4-10 Component-wise statistical analysis

Component	Method	Initial Value	Final Value	Convergence Rate	Variance ($\sigma^2 \times 10^3$)	Noise Sensitivity
R_{pos}	uPU	0.552 ± 0.032	0.087 ± 0.015	0.045 ± 0.008	9.2 ± 1.6	High
	Ramp-Loss	0.483 ± 0.018	0.038 ± 0.004	0.068 ± 0.006	1.3 ± 0.3	Low
R_{neg}	uPU	0.423 ± 0.028	0.063 ± 0.009	0.052 ± 0.009	7.5 ± 1.3	Very High
	Ramp-Loss	0.315 ± 0.015	0.028 ± 0.003	0.079 ± 0.007	0.9 ± 0.2	Moderate
Max Operator Activation	uPU	42% of epochs	18% of epochs	-	4.1 ± 0.9	-
	Ramp-Loss	15% of epochs	5% of epochs	-	0.8 ± 0.2	-

(Source: Created by author)

The component analysis reveals several key mechanisms:

- 1) Reduced Initial Risk: Bounded losses start with 12-25% lower initial risk values, providing a better starting point for optimization.
- 2) Faster Convergence: Both R_{pos} and R_{neg} components converge 35-50% faster with bounded losses, explaining the reduced training time.
- 3) Lower Variance: Component variances are reduced by 85-90%, explaining the smoother optimization trajectories.
- 4) Reduced Rectification Activation: The $\max(0, \cdot)$ operator activates significantly less frequently (15% vs 42% of epochs initially), indicating more stable risk estimates that rarely become negative.

The ablation studies collectively demonstrate that bounded loss functions provide robustness through multiple complementary mechanisms: direct outlier limitation, gradient stabilization, risk component balancing, and reduced rectification dependency. These mechanisms operate synergistically to produce the substantial performance improvements observed across all experimental conditions.

4.6 Discussion on Robustness and Generalization

4.6.1 Interpretation of Theoretical Bounds

The experimental results provide compelling empirical validation of the theoretical bounds established in Chapter III. The bounded risk distortion theorem (Theorem 1) predicted that label noise would impact performance proportionally to ηC (noise rate $\times$ loss bound). Our experimental observations align precisely with this prediction: as shown in Table 4.6, performance degradation scales linearly with noise rate, with bounded-loss methods exhibiting shallower slopes corresponding to their smaller effective C values.

The practical manifestation of the $2\pi\eta C$ bound can be quantified as follows: for the ramp loss ($C=1$), maximum risk distortion at $\eta=0.4$ is bounded by 0.8π , while for logistic loss (effectively unbounded), distortion can be unbounded. This explains the 19.5% versus 31.8% accuracy degradation observed empirically. The boundedness creates a protective envelope that limits the worst-case influence of corrupted labels, ensuring graceful degradation rather than catastrophic failure. Figure 4-8 shows the theoretical bound validation.




Figure 4-8 Theoretical bound validation (Source: Created by author)

The uniform convergence guarantee (Theorem 2) explains the dramatically better training stability observed in Table 4-7. The bounded gradient property (Definition 1) directly enables smoother optimization, with our methods achieving 74% lower loss variance than baselines. These theoretical foundations transform from mathematical abstractions to practical design principles, guiding the development of robust learning systems.

4.6.2 Practical Implications for Noisy PU Scenarios

The experimental findings yield actionable insights for practitioners facing label noise in real-world PU learning scenarios. Three key implications emerge:

1) Method Selection Guidelines: Based on comprehensive evaluation, we recommend:

- For $\eta < 0.2$: Standard nnPU may suffice
- For $0.2 \leq \eta \leq 0.4$: Truncated logistic loss ($\tau=1.5-2.0$)
- For $\eta > 0.4$: Ramp loss for maximum robustness
- When noise level unknown: Start with ramp loss for safety

2) Performance Expectations: Practitioners can now anticipate realistic performance bounds. For medical diagnosis with estimated $\eta=0.3$ label noise, our framework maintains 85-88% accuracy where traditional methods drop to 75-78%, potentially impacting clinical decision reliability.

3) Implementation Considerations: The 15% training time reduction, combined with dramatically improved stability, makes bounded-loss PU learning not only more accurate but also more computationally efficient and easier to deploy in production environments.

4.7 Chapter Summary

This chapter presents comprehensive experimental validation of the proposed bounded-loss PU learning framework across diverse datasets and noise conditions, empirically confirming theoretical claims and demonstrating practical advantages.

Section 4.1 details the experimental setup using three datasets: synthetic Gaussian (2,000 samples) for controlled analysis, UCI Heart (303 instances) for medical validation, and MNIST-PU (10,000 images) for high-dimensional testing. Noise is injected at rates $\eta \in \{0.0, 0.2, 0.4, 0.6\}$ by replacing clean positives with randomly selected negatives.

Section 4.2 presents performance results. Bounded-loss methods consistently outperform baselines, with advantages magnifying as noise increases. At $\eta=0.6$ on synthetic data, ramp-loss achieves 69.1% accuracy versus 56.3% for nnPU and 52.1% for uPU. On UCI Heart, ramp-loss maintains 73.5% accuracy compared to 62.8% for nnPU. Training stability improves dramatically, with loss variance reduced by 70-80%.

Section 4.3 compares methods directly. At $\eta=0.4$, ramp-loss achieves average improvements of 20.1% over uPU and 12.8% over nnPU. Performance degradation from clean to 60% noise is 19.5% for ramp-loss versus 31.8% for uPU, representing 38.7% reduction in degradation rate.

Section 4.4 analyzes training behavior. Bounded losses converge in half the epochs (38 vs 82) with 74% lower loss variance and 15% reduced training time. Optimal truncation threshold τ decreases with noise: $\tau=2.0$ for clean data, $\tau=1.5$ for moderate noise, $\tau=1.0$ for severe noise.

Section 4.5 presents ablation studies. Bounded loss alone provides 8.8% improvement; rectification alone provides 3.5%; their combination achieves 11.9%, demonstrating synergy. Component analysis shows 85-90% variance reduction and decreased rectification activation from 42% to 15% of epochs.

Section 4.6 discusses practical implications. Results validate theoretical bounds on risk distortion. Recommended guidelines: nnPU for $\eta < 0.2$, truncated logistic ($\tau=1.5-2.0$) for $0.2 \leq \eta \leq 0.4$, ramp loss for $\eta > 0.4$. For medical diagnosis with $\eta=0.3$, the framework maintains 85-88% accuracy versus 75-78% for traditional methods, demonstrating practical value in high-stakes applications.

Chapter V CONCLUSIONS

This research has systematically investigated the challenge of label noise in Positive-Unlabeled learning and developed a robust solution through the integration of bounded loss functions. Our comprehensive experimental evaluation across synthetic and real-world datasets has consistently demonstrated that bounded loss functions, specifically ramp loss and truncated logistic loss provide substantial improvements in both performance stability and classification accuracy under varying levels of label contamination. The proposed framework maintains the theoretical advantages of established PU learning methods while significantly enhancing their resilience to the practical reality of imperfect labels. Key empirical findings reveal that bounded losses reduce performance degradation under noise by approximately 40% compared to traditional approaches, while simultaneously improving training stability by reducing loss variance by 70-80% in high-noise conditions. These results validate our core hypothesis that boundedness in loss function design represents a principled and effective strategy for robust weakly supervised learning.

This thesis makes several significant contributions to the fields of Positive-Unlabeled learning and robust machine learning. First, we have developed a novel theoretical framework that formally characterizes the robustness properties of bounded loss functions in PU learning, establishing provable bounds on risk distortion under label noise. Second, we have introduced a practical algorithmic implementation that integrates bounded losses into the non-negative PU learning paradigm, creating a method that is both theoretically grounded and computationally efficient. Third, our extensive empirical validation across multiple domains provides compelling evidence for the effectiveness of bounded losses, establishing new performance benchmarks for noisy PU learning scenarios. Beyond these direct contributions, this work bridges previously disconnected research areas, weakly supervised learning, robust statistics, and loss function design demonstrating how insights from one domain can address fundamental challenges in another. The framework developed here represents not merely an incremental improvement but a conceptual advance in how we approach learning from imperfect data.

While this research has demonstrated significant advancements in robust PU learning, several inherent limitations must be acknowledged to properly contextualize its contributions and guide future work. First, the adopted noise model simplifies reality by assuming uniform random contamination; however, real-world label noise often exhibits more complex patterns such as instance-dependent ambiguity or structured corruption based on semantic similarity

between classes. Second, the current theoretical and algorithmic implementation focuses primarily on binary classification, leaving the essential extensions to multi-class and multi-label PU learning where the interplay of noise across multiple positive classes introduces additional complexity as important open challenges.

Third, our formal robustness analysis relies on certain regularity assumptions, including Lipschitz continuity and bounded loss gradients, which facilitate the derivation of guarantees but may not hold universally in all practical scenarios, particularly with highly non-convex, over-parameterized deep learning models. Fourth, while the measured computational benefits in training stability and convergence are substantial on conventional hardware, deployment on specialized edge devices or resource-constrained platforms may present additional challenges related to memory footprint and inference latency. Finally, the framework's core reliance on a reasonably accurate class prior (π) estimation represents a potential vulnerability; in settings of extreme noise or severe class imbalance, prior misspecification could interact deleteriously with label contamination in complex ways not fully addressed within the current theoretical bounds, warranting further investigation into joint robust estimation strategies.

The natural progression of this work involves a deeper and more systematic integration of bounded loss principles into contemporary deep learning architectures. While our current implementation demonstrates the benefits of bounded surrogates within a standard neural network framework, the full potential of bounded losses in synergy with powerful representation learning mechanisms remains largely unexplored. Future research should critically investigate how the intrinsic property of boundedness interacts with and influences feature learning dynamics, particularly in contexts where the architecture's inductive biases may either inadvertently amplify or actively mitigate the effects of label noise. A dedicated inquiry into how gradient clipping at the loss level propagates through layer-wise representations could yield new insights into training stability and generalization under corruption.

Furthermore, exploring specialized bounded loss variants specifically engineered for advanced architectural paradigms such as convolutional networks for vision, recurrent models for sequences, and transformer-based structures for language could unlock robust PU learning in high-dimensional and structured data domains. This architectural tailoring would address the unique optimization landscapes and noise sensitivities of each model family. Ultimately, the most promising avenue may lie in the synergistic combination of bounded loss functions with modern regularization techniques, adaptive attention mechanisms, and other architectural innovations, paving the way for a new generation of deeply robust, end-to-end trainable PU

learning systems that are resilient by design across diverse data modalities and application scenarios.

Moving beyond uniform noise assumptions represents a critical next step. Real-world label noise often exhibits systematic patterns: certain instance types are more prone to mislabeling, noise rates may vary across data subsets, and contamination may correlate with specific features. Future work should develop adaptive bounded loss frameworks that dynamically adjust their robustness parameters based on estimated noise characteristics. This could involve learning instance-specific loss bounds, incorporating meta-learning to adapt to noise patterns, or developing hierarchical models that distinguish between different noise regimes within the same dataset. Such approaches would better capture the complexity of real-world data imperfections while maintaining the theoretical guarantees established in this work.

The principles developed in this thesis naturally extend to broader weakly supervised learning paradigms. Future research should explore the integration of bounded loss functions with semi-supervised learning techniques, where limited labeled data (potentially noisy) is supplemented with larger unlabeled collections. Similarly, combining bounded losses with self-supervised pretext tasks could create more robust representation learning frameworks that are less susceptible to label quality issues. The synergy between boundedness for robustness and modern unsupervised learning for representation could yield transformative advances in learning from partially and imperfectly labeled data. Additionally, exploring connections with related paradigms such as noisy label learning, partial label learning, and learning with label proportions could yield valuable cross-pollination of ideas and techniques.

This thesis has addressed one of the most persistent challenges in practical machine learning: how to learn effectively when both the quantity and quality of labels are limited. By developing and validating a bounded loss framework for Positive-Unlabeled learning, we have advanced the state of the art in robust weakly supervised learning while providing both theoretical insights and practical tools for real-world application. The demonstrated improvements in noise resilience, training stability, and generalization performance represent meaningful progress toward machine learning systems that can operate reliably in the imperfect data environments that characterize most practical applications.

References

- [1] Roy P. Enhancing Real-World Robustness in AI: Challenges and Solutions[J]. JOURNAL OF RECENT TRENDS IN COMPUTER SCIENCE AND ENGINEERING, 2024, 12: 34-49.
- [2] Zulfiqar M I, Khalid A, Chen L, et al. Uncertainty-Aware Neighbor Calibration for Positive and Unlabeled Learning in Large Machine Learning Models[J]. IEEE Trans Emerg Top Comput Intell, 2026, 10: 543-556.
- [3] Wei J, Wu Y, Shi B, et al. ALM-PU: positive and unlabeled learning with constrained optimization[J]. Mach Learn, 2025, 114: 210.
- [4] Bekker J, Davis J. Learning from positive and unlabeled data: a survey[J]. Mach Learn, 2020, 109: 719-760.
- [5] Fredriksson T, Mattos D I, Bosch J, et al. An Empirical Evaluation of Algorithms for Data Labeling[C]//2021 IEEE 45th Annual Computers, Software, and Applications Conference (COMPSAC). IEEE, 2021: 201-209.
- [6] Vranckx I, Raymaekers J, De Ketelaere B, et al. Real-time discriminant analysis in the presence of label and measurement noise[J]. Chemometrics and Intelligent Laboratory Systems, 2021, 208: 104197.
- [7] Jaskie K, Spanias A. Positive And Unlabeled Learning Algorithms And Applications: A Survey[C]//2019 10th International Conference on Information, Intelligence, Systems and Applications (IISA). IEEE, 2019: 1-8.
- [8] Saha S, Hassan A, Paik J H. A Positive-Unlabeled Learning Approach with Self-correcting Regularized Risk[J]. IEEE Transactions on Artificial Intelligence, 2025: 1-13.
- [9] Wang Q, Ma Y, Zhao K, et al. A Comprehensive Survey of Loss Functions in Machine Learning[J]. Annals of Data Science, 2022, 9: 187-212.
- [10] Qaraei M, Schultheis E, Gupta P, et al. Convex Surrogates for Unbiased Loss Functions in Extreme Classification With Missing Labels[C]//Proceedings of the Web Conference 2021. New York, NY, USA: ACM, 2021: 3711-3720.
- [11] Saunders J D, Freitas A A. Evaluating the Predictive Performance of Positive-Unlabelled Classifiers[J]. ACM SIGKDD Explorations Newsletter, 2022, 24: 5-11.
- [12] Wang Y, Chen Q, Zhang H, et al. Biomedical Information Retrieval with Positive-Unlabeled Learning and Knowledge Graphs[J]. ACM Trans Intell Syst Technol, 2024.

- [13] Wolf D, Regnery S, Tarnawski R, et al. Weakly Supervised Learning with Positive and Unlabeled Data for Automatic Brain Tumor Segmentation[J]. *Applied Sciences*, 2022, 12: 10763.
- [14] Deng S, Han B, Li X, et al. Positive and unlabeled learning on generating strategy for weakly anomaly detection[J]. *Signal Image Video Process*, 2025, 19: 171.
- [15] Bekker J, Robberechts P, Davis J. Beyond the Selected Completely at Random Assumption for Learning from Positive and Unlabeled Data[M]//*Lecture Notes in Computer Science*. 2020: 71-85.
- [16] Kwon Y, Kim W, Sugiyama M, et al. Principled analytic classifier for positive-unlabeled learning via weighted integral probability metric[J]. *Mach Learn*, 2020, 109: 513-532.
- [17] Awasthi L, Danso E. Robust Positive-Unlabeled Learning via Bounded Loss Functions under Label Noise[J]. *Scientific Journal of Engineering Research*, 2025, 1: 140-152.
- [18] DU PLESSIS M C, SUGIYAMA M. Class Prior Estimation from Positive and Unlabeled Data[J]. *IEICE Trans Inf Syst*, 2014, E97.D: 1358-1362.
- [19] M S V, Yuan S, Wu X. Fraud Detection via Contrastive Positive Unlabeled Learning[C]//*2022 IEEE International Conference on Big Data (Big Data)*. IEEE, 2022: 1475-1484.
- [20] Song K, Liu C, Jiang D. A positive-unlabeled learning approach for industrial anomaly detection based on self-adaptive training[J]. *Neurocomputing*, 2025, 647: 130488.
- [21] Zhu Z, Wang L, Zhao P, et al. Robust Positive-Unlabeled Learning via Noise Negative Sample Self-correction[C]//*Proceedings of the 29th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*. New York, NY, USA: ACM, 2023: 3663-3673.
- [22] Johnson J M, Khoshgoftaar T M. A Survey on Classifying Big Data with Label Noise[J]. *Journal of Data and Information Quality*, 2022, 14: 1-43.
- [23] Wang X, Zhu R, Xue J-H. UC-PUAL: A universally consistent classifier of positive-unlabelled data[J]. *Pattern Recognit*, 2026, 169: 111892.
- [24] Tang X, Xu C, Tao H, et al. Confidence-Based PU Learning With Instance-Dependent Label Noise[J]. *IEEE Trans Neural Netw Learn Syst*, 2025, 36: 14283-14297.
- [25] Li Y-F, Guo L-Z, Zhou Z-H. Towards Safe Weakly Supervised Learning[J]. *IEEE Trans Pattern Anal Mach Intell*, 2019: 1-1.
- [26] Teisseyre P, Furmańczyk K, Mielniczuk J. Verifying the Selected Completely at Random Assumption in Positive-Unlabeled Learning[J/OL]. *arXiv preprint arXiv:2404.00145*, 2024.

- [27] Zhao Y, Zhang M, Zhang C, et al. A boosting framework for positive-unlabeled learning[J]. *Stat Comput*, 2025, 35: 2.
- [28] Kiryo R, Niu G, Du Plessis M C, et al. Positive-Unlabeled Learning with Non-Negative Risk Estimator[C]//*Adv Neural Inf Process Syst*. 2017.
- [29] Garg S, Wu Y, Smola A, et al. Mixture Proportion Estimation and PU Learning: A Modern Approach[J/OL]. *Adv Neural Inf Process Syst*, 2021.
- [30] Han Z, He R, Li T, et al. Semi-Supervised Screening of COVID-19 from Positive and Unlabeled Data with Constraint Non-Negative Risk Estimator[C]//*Information Processing in Medical Imaging*. 2021: 611-623.
- [31] Hummel H I, Gansekoele A, Bhulai S, et al. The computation of generalized embeddings for underwater acoustic target recognition using contrastive learning[J]. *Applied Acoustics*, 2026, 243: 111103.
- [32] Li J, Ma X, Shi Y. Dual-Stage Clean-Sample Selection for Incremental Noisy Label Learning[J]. *Bioengineering*, 2025, 12: 743.
- [33] Song H, Kim M, Park D, et al. Learning From Noisy Labels With Deep Neural Networks: A Survey[J]. *IEEE Trans Neural Netw Learn Syst*, 2023, 34: 8135-8153.
- [34] Zavitsanos E, Bougiatiotis K, Sideras A, et al. Positive-Unlabeled Learning for Financial Misstatement Detection under Realistic Constraints[C]//*Proceedings of the 6th ACM International Conference on AI in Finance*. New York, NY, USA: ACM, 2025: 864-872.
- [35] Nguyen T V, Diakiw S M, VerMilyea M D, et al. Efficient automated error detection in medical data using deep-learning and label-clustering[J]. *Sci Rep*, 2023, 13.
- [36] Akujuobi U, Chen J, Elhoseiny M, et al. Temporal Positive-unlabeled Learning for Biomedical Hypothesis Generation via Risk Estimation[C]//*Adv Neural Inf Process Syst*. 2020: 4597-4609.
- [37] Naser M Z. From failure to fusion: A survey on learning from bad machine learning models[J]. *Information Fusion*, 2025, 120: 103122.
- [38] Sevetlidis V, Pavlidis G, Mouroutsos S G, et al. Dense-PU: Learning a Density-Based Boundary for Positive and Unlabeled Learning[J]. *IEEE Access*, 2024, 12: 90287-90298.
- [39] Hajiyevev A, Xu J. High-Dimensional Regression Analysis and Artificial Intelligence[M]. Singapore: Springer Nature Singapore, 2026.
- [40] Leqi L, Huang A, Lipton Z C, et al. Supervised Learning with General Risk Functionals[C]//*International Conference on Machine Learning*. 2022.

- [41] Beckley J. Advanced Risk Assessment Techniques: Merging Data-Driven Analytics with Expert Insights to Navigate Uncertain Decision-Making Processes[J]. *International Journal of Research Publication and Reviews*, 2025, 6: 1454-1471.
- [42] Cao Y, Xiang H, Zhang H, et al. Anomaly Detection Based on Isolation Mechanisms: A Survey[J]. *Machine Intelligence Research*, 2025, 22: 849-865.
- [43] Toleu A, Tolegen G, Pak A, et al. Noise-Aware Direct Preference Optimization for RLAIIF[J]. *Applied Sciences*, 2025, 15: 10328.
- [44] Hasan M, Shifa S I, Niaz K, et al. Continuous Data Curation and Valuation for Long-Term Machine Learning Model Health: A Comprehensive Review[J]. *EJSMT*, 2026, 2: 58-78.
- [45] Paleyes A, Urma R-G, Lawrence N D. Challenges in Deploying Machine Learning: A Survey of Case Studies[J]. *ACM Comput Surv*, 2023, 55: 1-29.
- [46] Frenay B, Verleysen M. Classification in the Presence of Label Noise: A Survey[J]. *IEEE Trans Neural Netw Learn Syst*, 2014, 25: 845-869.
- [47] Wen X-C, Wang X, Gao C, et al. When Less is Enough: Positive and Unlabeled Learning Model for Vulnerability Detection[C]//2023 38th IEEE/ACM International Conference on Automated Software Engineering (ASE). IEEE, 2023: 345-357.
- [48] Chou Y-T, Niu G, Lin H-T, et al. Unbiased Risk Estimators Can Mislead: A Case Study of Learning with Complementary Labels[C]//International conference on machine learning. 2020: 1929-1938.
- [49] Algan G, Ulusoy İ. Label Noise Types and Their Effects on Deep Learning[J/OL]. *arXiv preprint arXiv:2003.10471*, 2020.
- [50] Wang C, Pu J, Xu Z, et al. Asymmetric Loss for Positive-Unlabeled Learning[C]//2021 IEEE International Conference on Multimedia and Expo (ICME). IEEE, 2021: 1-6.
- [51] Liao Z, Hu S, Xie Y, et al. Instance-dependent Label Distribution Estimation for Learning with Label Noise[J]. *Int J Comput Vis*, 2025, 133: 2568-2580.
- [52] Wang Q, Han B, Liu T, et al. Tackling Instance-Dependent Label Noise via a Universal Probabilistic Model[C]//Proceedings of the AAAI Conference on Artificial Intelligence. 2021, 35: 10183-10191.
- [53] Song B, Zhao S, Dang L, et al. A survey on learning from data with label noise via deep neural networks[J]. *Systems Science & Control Engineering*, 2025, 13.
- [54] Zhang Y, Wei R, Wei S, et al. Accurate coal ignition temperature prediction through interpretable deep transfer learning with small-sample feature-level adaptive noise augmentation[J]. *Energy Convers Manag*, 2026, 348: 120643.

- [55] Zhang S, Xie L, Chen Y, et al. Inter-class margin climbing with cost-sensitive learning in neural network classification[J]. Knowl Inf Syst, 2025, 67: 1993-2016.
- [56] Zhou C, Meng H, Li M, et al. On Learning Label Noise Robust Networks via Regularization: A Topological View[J]. IEEE Trans Neural Netw Learn Syst, 2025, 36: 16156-16170.
- [57] Ghosh A, Kumar H, Sastry P S. Robust Loss Functions under Label Noise for Deep Neural Networks[C]//Proceedings of the AAAI Conference on Artificial Intelligence. 2017, 31.
- [58] Terven J, Cordova-Esparza D-M, Romero-González J-A, et al. A comprehensive survey of loss functions and metrics in deep learning[J]. Artif Intell Rev, 2025, 58: 195.
- [59] Ding K, Shu J, Meng D, et al. Improve Noise Tolerance of Robust Loss via Noise-Awareness[J]. IEEE Trans Neural Netw Learn Syst, 2025, 36: 13189-13203.
- [60] Zhou X, Liu X, Zhai D, et al. Asymmetric Loss Functions for Noise-Tolerant Learning: Theory and Applications[J]. IEEE Trans Pattern Anal Mach Intell, 2023, 45: 8094-8109.
- [61] Li C, Liu K, Liu S. A Survey of Loss Functions in Deep Learning[J]. Mathematics, 2025, 13: 2417.
- [62] Xu Y, Zhang D, Stefanidis A. Deep Neural Network Loss Landscapes: From Geometry to Generalization[J/OL]. Authorea, 2026.
- [63] Oleghe O. A predictive noise correction methodology for manufacturing process datasets[J]. J Big Data, 2020, 7: 89.
- [64] Xu C, Song P, Tang S, et al. Alleviating Confirmation Bias in Learning with Noisy Labels via Two-Network Collaboration[J]. ACM Trans Intell Syst Technol, 2025, 16: 1-21.
- [65] Jiang L, Zhou Z, Leung T, et al. MentorNet: Learning Data-Driven Curriculum for Very Deep Neural Networks on Corrupted Labels[C]//International conference on machine learning. 2018: 2304-2313.
- [66] Wang X, Chen Y, Zhu W. A Survey on Curriculum Learning[J]. IEEE Trans Pattern Anal Mach Intell, 2021: 1-1.
- [67] Kirtas M, Passalis N, Mourgiyas-Alexandris G, et al. Robust Architecture-Agnostic and Noise Resilient Training of Photonic Deep Learning Models[J]. IEEE Trans Emerg Top Comput Intell, 2023, 7: 140-149.
- [68] Bai Y, Yang E, Han B, et al. Understanding and Improving Early Stopping for Learning with Noisy Labels[C]//Adv Neural Inf Process Syst. 2021: 24392-24403.
- [69] Nagarajan V, Kolter J Z. Uniform convergence may be unable to explain generalization in deep learning[C]//Adv Neural Inf Process Syst. 2019.

- [70] Dwivedi P, Islam B, Kajal M. Smooth gradient loss: a loss function for gradient regularization in deep learning optimization[J]. J Supercomput, 2025, 81: 1457.
- [71] Khan A, Ali A, Islam N, et al. Robust Extreme Learning Machine Using New Activation and Loss Functions Based on M-Estimation for Regression and Classification[J]. Sci Program, 2022, 2022: 1-10.
- [72] Shafique M, Naseer M, Theocharides T, et al. Robust Machine Learning Systems: Challenges, Current Trends, Perspectives, and the Road Ahead[J]. IEEE Des Test, 2020, 37: 30-57.
- [73] Tian Y, Mirzabagheri M, Bamakan S M H, et al. Ramp loss one-class support vector machine; A robust and effective approach to anomaly detection problems[J]. Neurocomputing, 2018, 310: 223-235.
- [74] Hasan M K, Pal C. A New Smooth Approximation to the Zero One Loss with a Probabilistic Interpretation[J]. ACM Trans Knowl Discov Data, 2020, 14: 1-28.
- [75] Chen H, Guo C, Xiong H, et al. Sparse additive machine with ramp loss[J]. Analysis and Applications, 2021, 19: 509-528.
- [76] Chauve C, Ponty Y, Wallner M. Counting and sampling gene family evolutionary histories in the duplication-loss and duplication-loss-transfer models[J]. J Math Biol, 2020, 80: 1353-1388.
- [77] Ju L, Wang X, Wang L, et al. Improving Medical Images Classification With Label Noise Using Dual-Uncertainty Estimation[J]. IEEE Trans Med Imaging, 2022, 41: 1533-1546.
- [78] Liu C, Dong Y, Xiang W, et al. A Comprehensive Study on Robustness of Image Classification Models: Benchmarking and Rethinking[J]. Int J Comput Vis, 2025, 133: 567-589.
- [79] Ran L, Li Y, Liang G, et al. Pseudo Labeling Methods for Semi-Supervised Semantic Segmentation: A Review and Future Perspectives[J]. IEEE Transactions on Circuits and Systems for Video Technology, 2025, 35: 3054-3080.
- [80] Li C, Huang C, Chen X, et al. Adversarial Constraints Guide Multi-View Learning for Semi-Supervised Change Detection[J]. IEEE Transactions on Geoscience and Remote Sensing, 2026: 1-1.
- [81] Joy U G, Rahman M M. AdaptiveSSL: a unified semi-supervised learning framework for robust classification via adaptive uncertainty calibration and dynamic labeling[J]. Knowl Inf Syst, 2026, 68: 41.
- [82] Silva Filho T, Song H, Perello-Nieto M, et al. Classifier calibration: a survey on how to assess and improve predicted class probabilities[J]. Mach Learn, 2023, 112: 3211-3260.

Research Output Achieved During Master's Degree Studies

I. Academic Papers

(Including published papers, papers accepted for publication, manuscripts under review, and manuscripts prepared or planned for submission that are directly related to the thesis)

No.	Author(s) (in order)	Paper Title	Journal/Level	Volume / Issue / Year / Pages	Related Chapter / Section of the Thesis	Indexing Status
1	Lalit Awasthi, Eric Danso	Robust Positive-Unlabeled Learning via Bounded Loss Functions under Label Noise	Scientific Journal of Engineering Research	vol. 1, no. 3, 2025	Chapter Three and Four	EI

Acknowledgement

The completion of this thesis marks a significant academic and personal achievement, one that would not have been possible without the generous support, profound guidance, and steadfast encouragement of numerous individuals and institutions to whom I extend my deepest gratitude.

My foremost and most sincere appreciation is reserved for my supervisor, Professor Tang Long, whose exceptional mentorship was the foundation of this entire endeavor. From the initial exploration of the research problem to the final refinement of this manuscript, his expert guidance, insightful critiques, and unwavering confidence in my work were indispensable. His patience in discussing theoretical nuances and his steadfast support through experimental challenges were pivotal in shaping this research and my growth as a researcher.

I extend my heartfelt thanks to my collaborator and friend, whose partnership was instrumental in this work. The countless hours spent in discussion, debugging complex code, analyzing experimental results, and navigating research obstacles were powered by our shared commitment. His technical expertise, logical clarity, and persistent optimism were a constant source of strength and inspiration, making this accomplishment a shared endeavor.

I am profoundly grateful to the faculty and staff of the School of Artificial Intelligence at Nanjing University of Information Science and Technology (NUIST). The rigorous academic training, access to critical computational resources, and the stimulating intellectual climate provided the essential bedrock for this research.

I acknowledge my debt to the broader scholarly communities in machine learning, weakly supervised learning, and robust statistics. The pioneering work of countless researchers, as well as the constructive feedback from peer reviewers on related publications, has profoundly influenced and strengthened the quality of this thesis.

My most profound and personal gratitude is for my family. To my parents, whose immeasurable sacrifices, boundless faith in the value of education, and unconditional love have been my unwavering anchor and inspiration throughout my life and studies.

Finally, I recognize the personal transformation inherent in the research process. This journey has instilled in me a deeper resilience, a more refined capacity for critical inquiry, and a persevering spirit. It has equipped me not only with specific knowledge but with a mindset poised for future challenges in scientific research. To everyone who has been a part of this pivotal chapter, thank you.